

Code Generation in MLIR

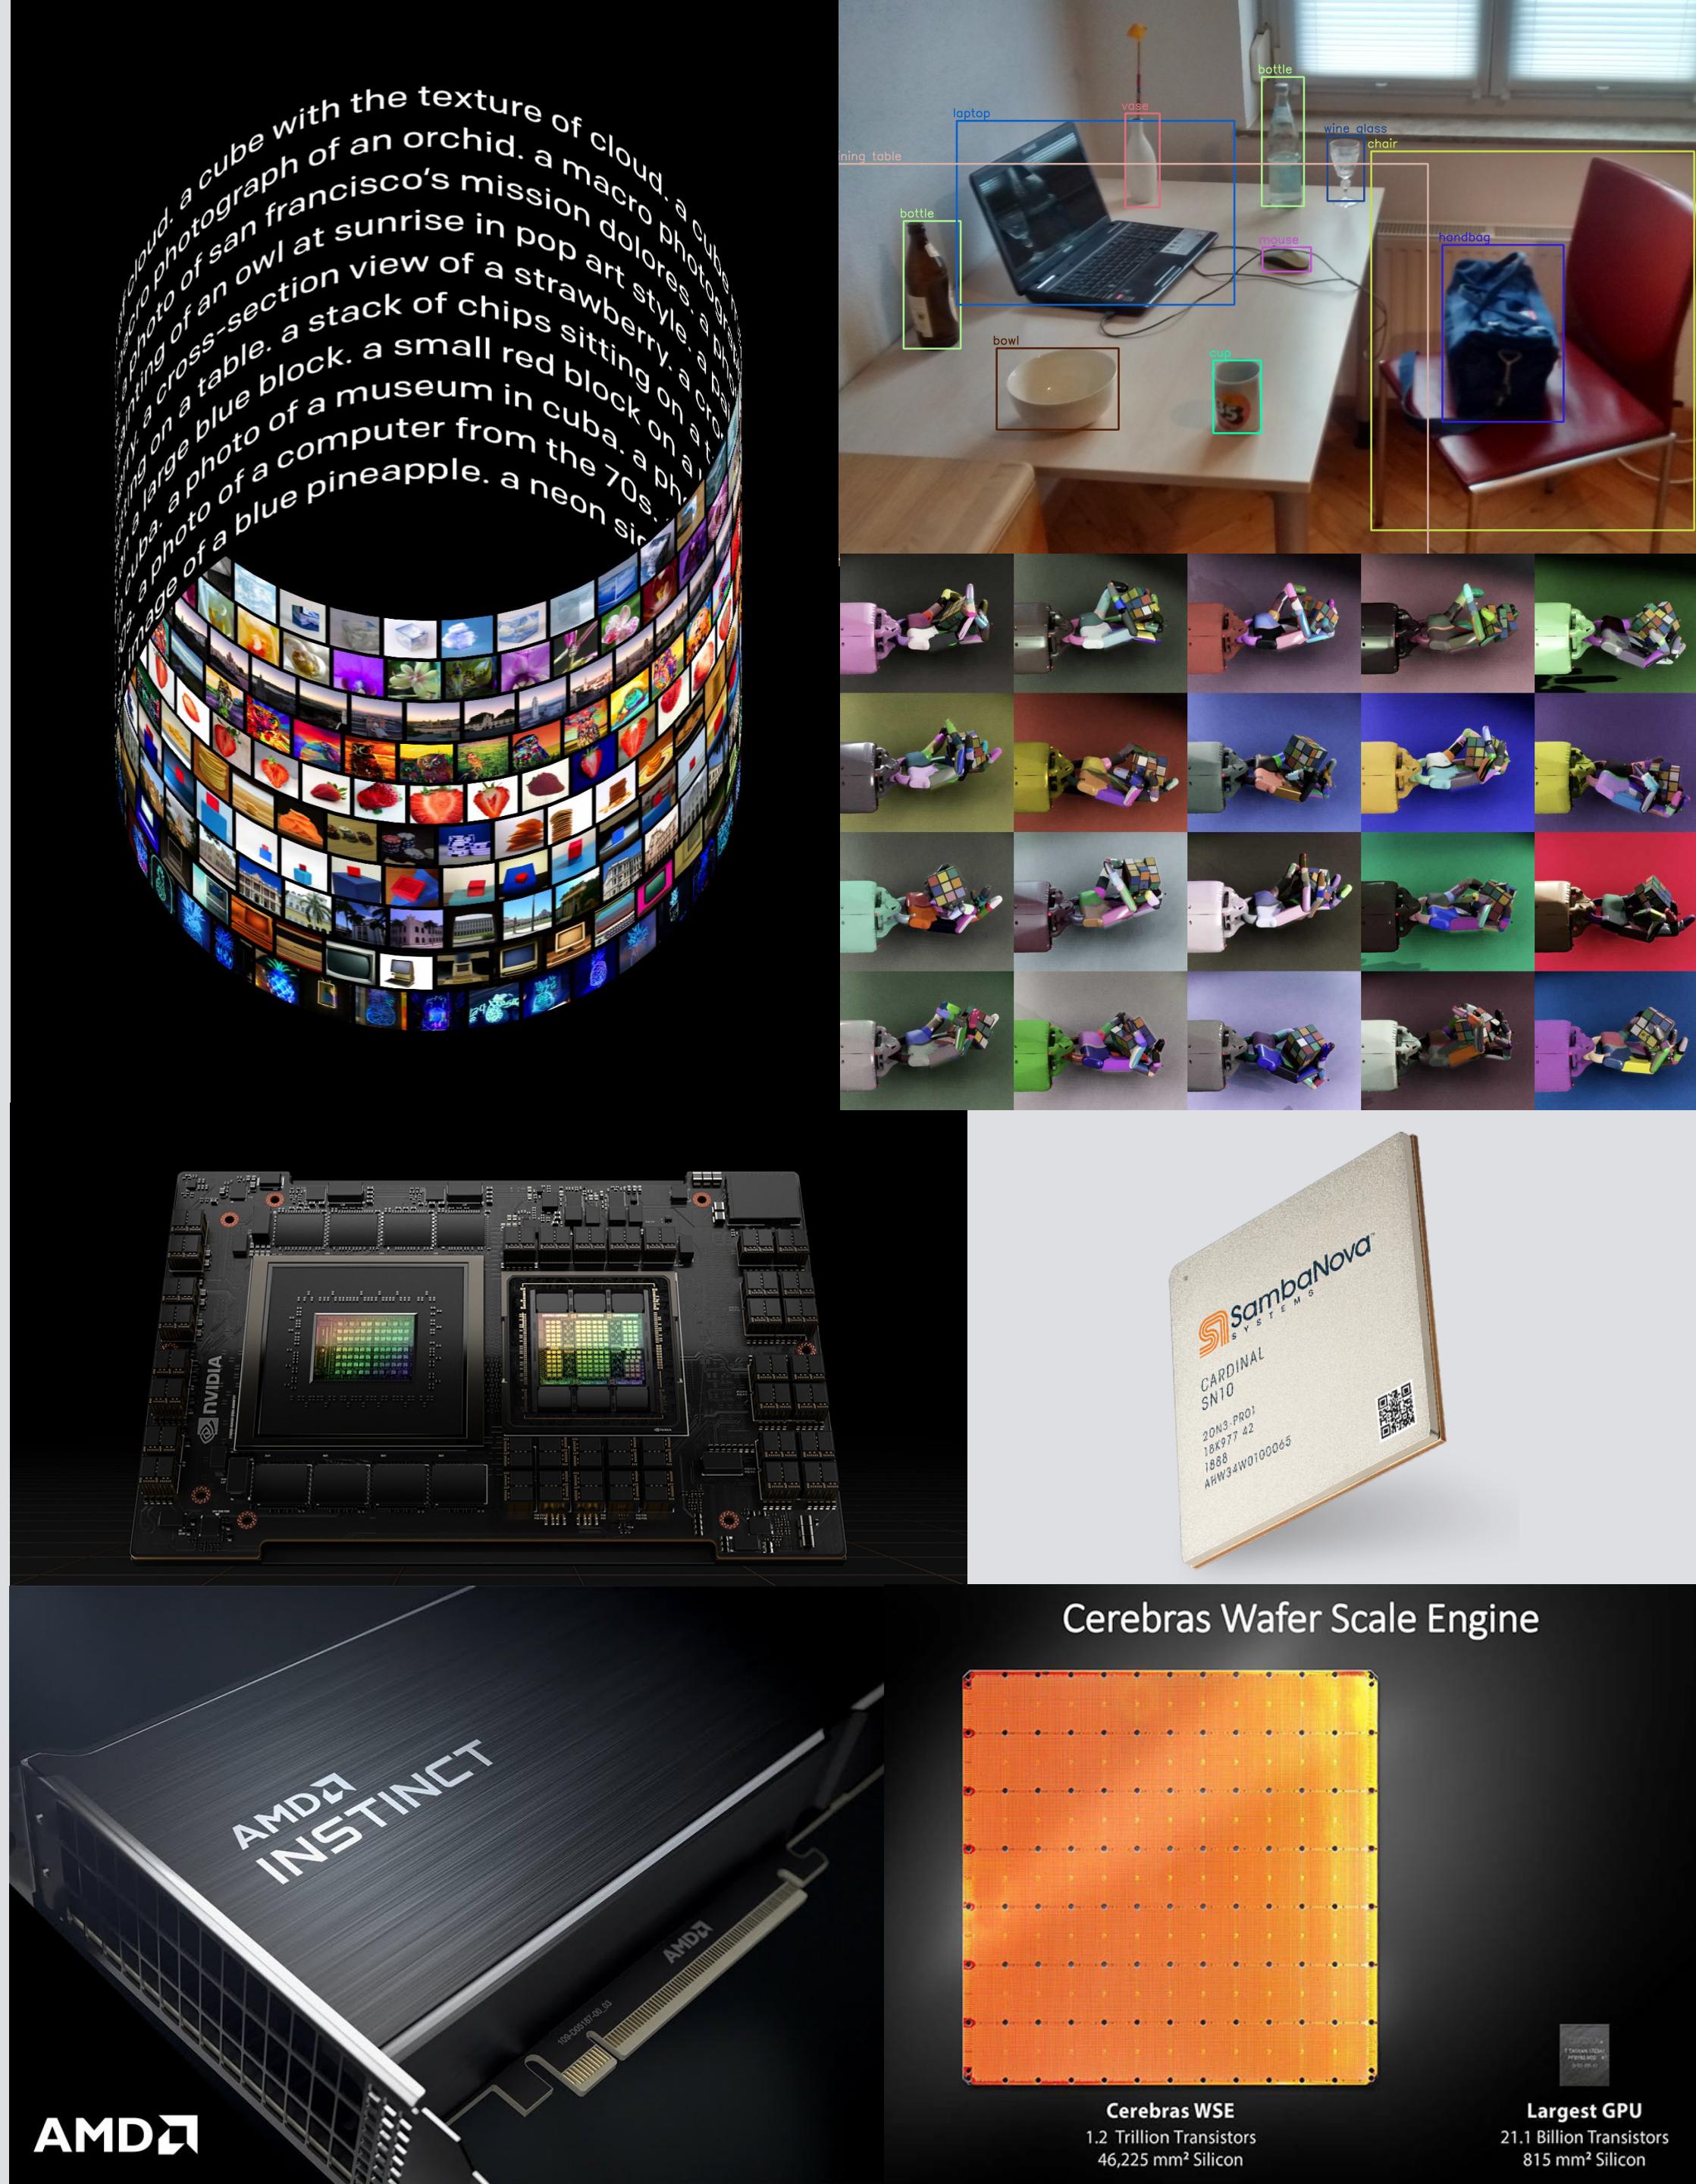
Harsh Menon (nod.ai)

Overview

- Motivation
- Code generation in LLVM/MLIR
- Dialects (Linalg, Vector, GPU, etc.)
- Walkthrough of code generation using IREE / SHARK
 - CPU and GPU code generation
 - Targeting custom accelerators
 - Auto-tuning
- Conclusion
- Acknowledgements
- References

Motivation

- Deep learning has become extremely pervasive spanning domains ranging from autonomous cars, natural language processing to medical imaging
- Large models are achieving state of the art results (such as transformers in NLP)
- Unfortunately, these large models take many months and millions of dollars to train on existing hardware
- On edge-based systems, inference dominates with latency being a key metric
- New hardware vendors have risen to the occasion with custom accelerators that address some of these concerns
- But as the number of models and hardware combinations explode, need a strong compiler infrastructure that can provide performance gains and is easily re-targetable to new hardware



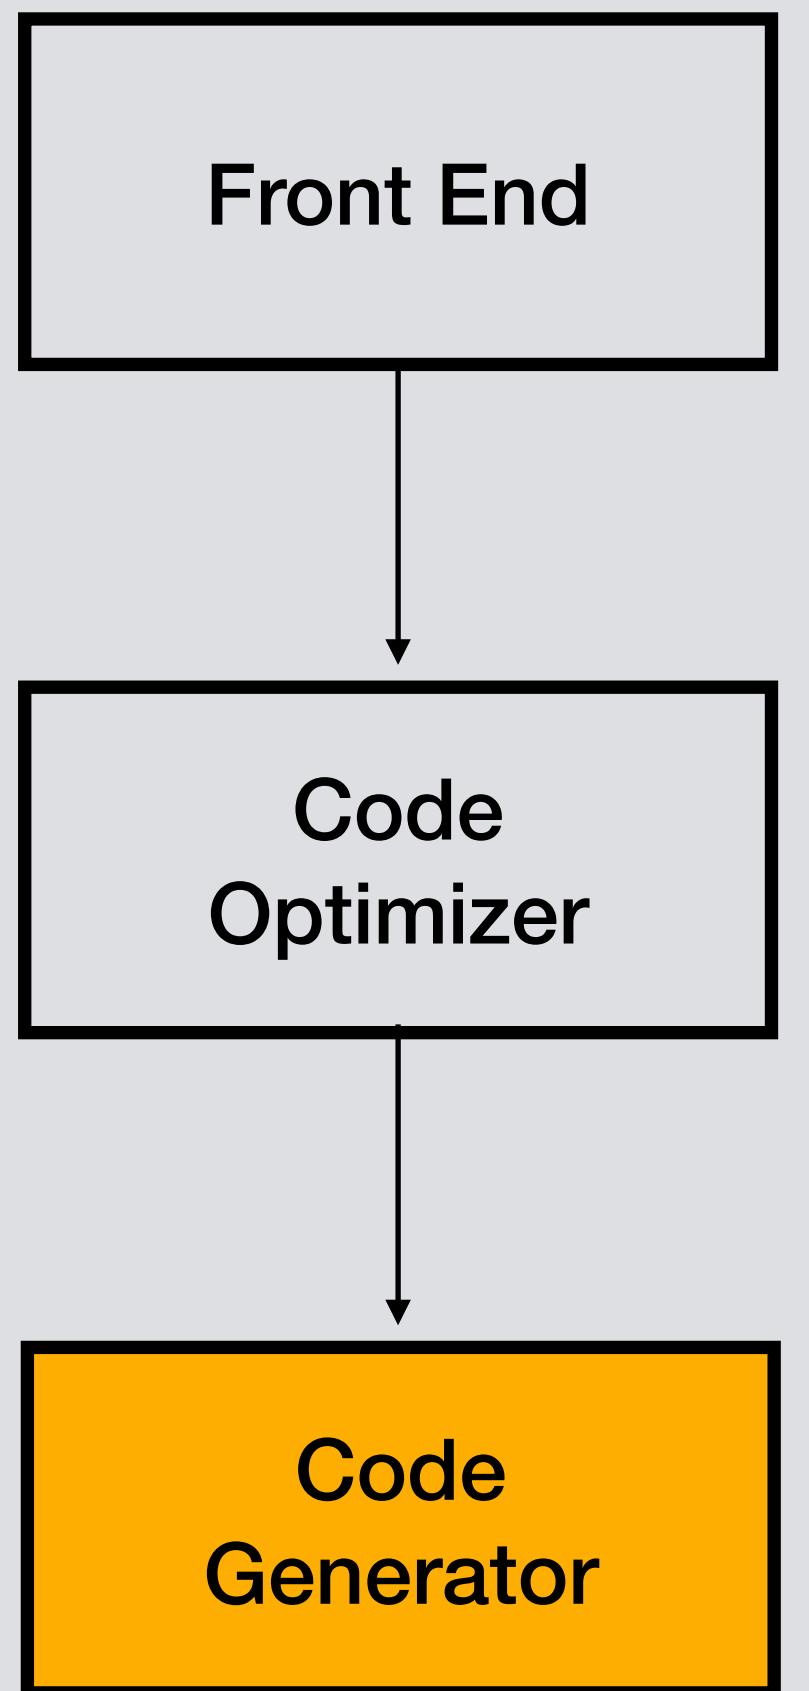
Motivation

- LLVM provides a strong compiler infrastructure that already scales to various hardware targets and can be used to fill the gap
- But LLVM IR is too low-level and many opportunities for optimization are missed if we start at that level of abstraction
- MLIR provides compiler infrastructure to handle varying levels of abstraction and provides a way to progressively lower to LLVM IR, leveraging best of both worlds
- In machine learning (ML), neural networks are defined in Python-based frameworks such as Tensorflow, PyTorch and JAX
- MLIR helps progressively lower computation graphs from their pythonic representation to LLVM IR



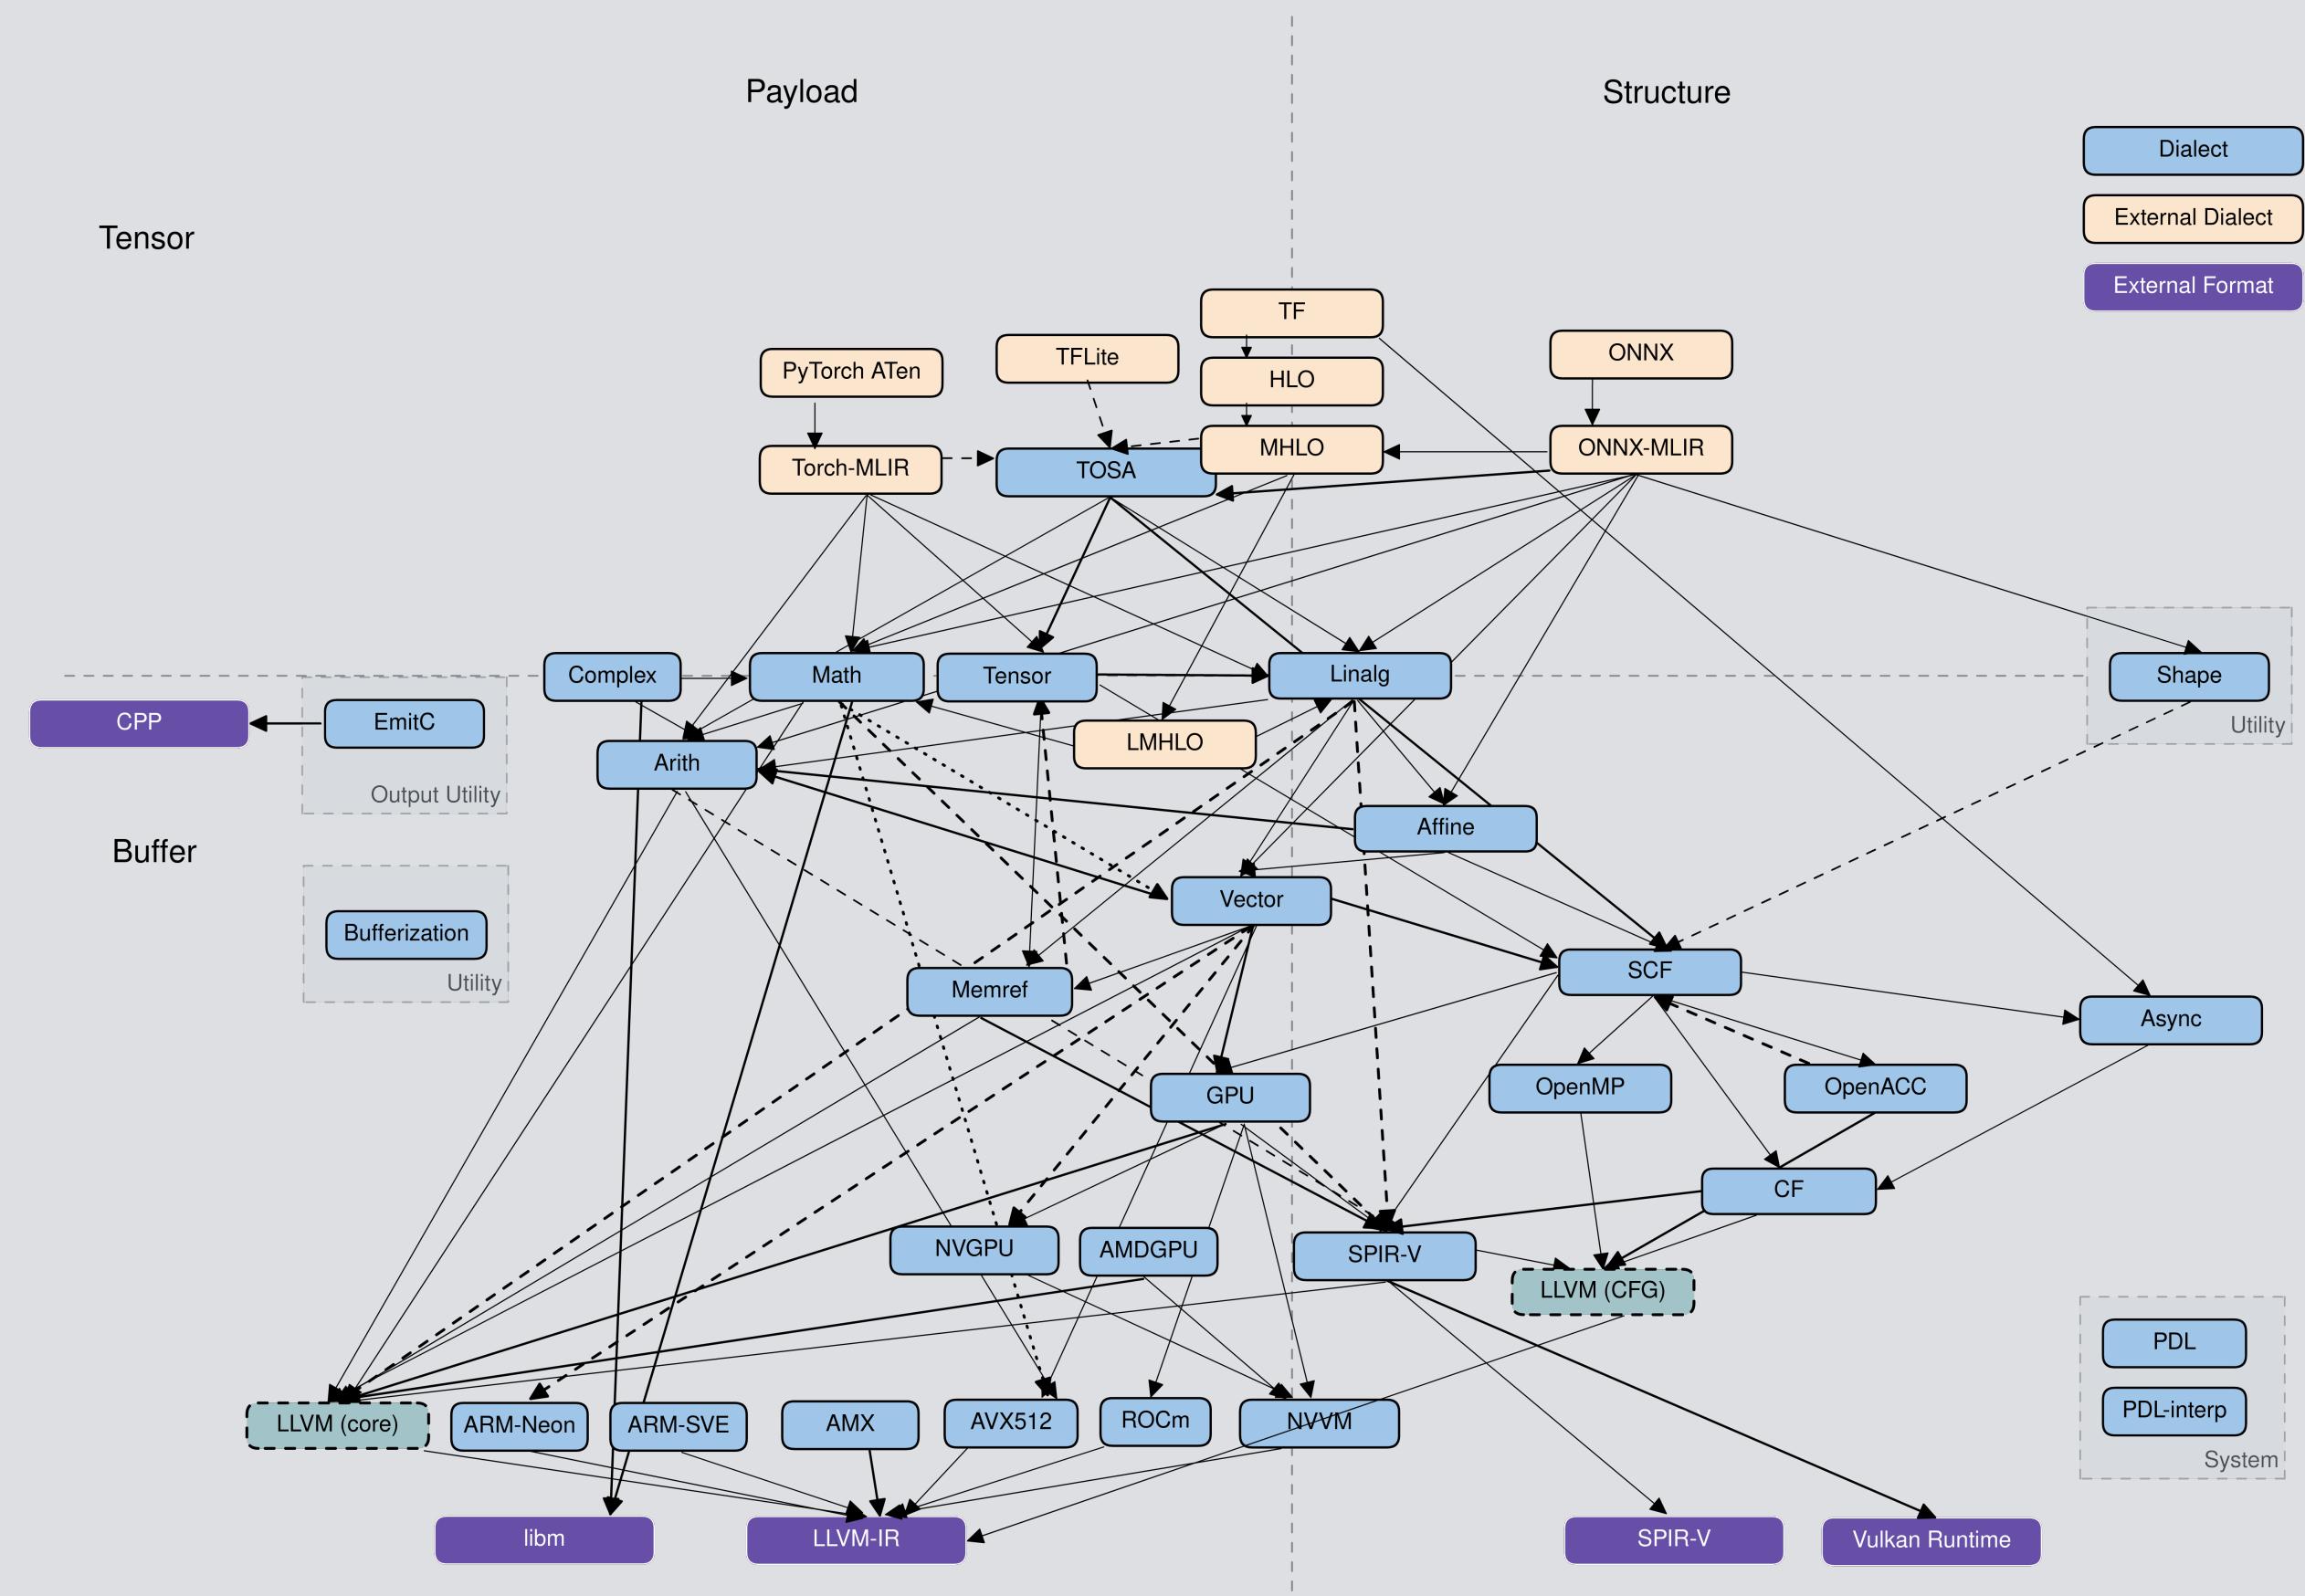
Code generation in LLVM

- Final phase in compiler pipeline
- Must make effective use of available resources while preserving program semantics
- Deals with problems such as instruction selection, register allocation and assignment, instruction ordering
- LLVM has several different backends such as X86, NVPTX, RISC-V etc. which contain the specific hardware definition (in terms of hardware instructions, registers etc. expressed in tablegen)
- During code generation pipeline, LLVM IR is lowered to the SelectionDAG, then MachineInstr, MCInst and finally compiled to bitcode
- MLIR leverages LLVM for this part of code generation
- Code generation in MLIR operates at a higher-level of abstraction and attempts to provide the missing infrastructure between high level ML programs and LLVM IR



Code generation in MLIR

- Figure on left shows dialect ecosystem in MLIR
- Starting with dialects on top (that closely follow the native ML frameworks), there are many paths to LLVM dialect
- Each dialect progressively lowers the abstraction going from tensors (immutable SSA) to memrefs
- Perform high-level optimizations at upper dialects, more hardware-specific optimizations as we approach LLVM dialect
- Once we arrive at LLVM dialect, we can translate to LLVM IR using mlir-translate and using LLVM to generate the final binary



Linalg Dialect

- Linalg dialect is used to represent perfectly nested loop computations making it easy to perform transformations like fusion, tiling, and loop interchange
- Operates on both tensors and memrefs
- Can be lowered to loops or affine expressions with computation in loop body
- Incorporates learnings from Halide, TVM, Tensor Comprehensions, XLA etc.
- Linalg defines a small set of core named ops (such as matmul, conv, pooling etc.) that the front-end dialects can lower to
- This maps the large set of operations to a smaller set of operations that the compiler can focus on optimizing
- The core workhorse of this dialect is the linalg.generic op

Linalg GenericOp

- **Indexing Maps**

- Capture data access patterns for each operand
- Domain represents point in iteration space, range represents a point in the operand's data space

- **Iterator Types**

- Specifies data dependence between iterations of the loop

- **Inputs and Outputs**

- Tied output and result operands

- **Compute Payload**

- Computation performed at each point in the iteration space
- Yielded value is result of computation
- Arguments are obtained from operands using indexing maps

```
%6 = linalg.generic {
    indexing_maps =
        [affine_map<(d0, d1, d2, d3) -> (d0, d1, d2, d3),
         affine_map<(d0, d1, d2, d3) -> (d0, d1)],
    iterator_types = ["parallel", "parallel", "reduction", "reduction"]
    ins(%1 : tensor<?x?x?x?xf32>) outs(%5: tensor<?x?xf32>) {
        ^bb0(%arg1: f32, %arg2: f32):
            %17 = arith.addf %arg2, %arg1 : f32
            linalg.yield %17 : f3
    } -> tensor<?x?xf32>
```

```
for (int i = 0; i < M; i++) {
    for (int j = 0; j < N; j++) {
        for (int k = 0; k < K; k++) {
            for (int l = 0; l < P; l++) {
                c[i][j] = c[i][j] + a[i][j][k][l];
            }
        }
    }
}
```

Vector Dialect

- Provides generic retargetable n-D vector abstractions
- Operations in vector dialect can progressively decompose to lower rank variants
- Can lower to LLVM instructions or directly target hardware intrinsics (mma_compute)
- Examples:
 - `vector.transfer_read/write` : bridge the gap between memory and vectors and contain enough information to encode read/write patterns such as broadcasted, permuted and masked accesses
 - `vector.outerproduct` : outer product operation (typically obtained from lowering of matmul) that can further be lowered to llvm fused-multiply add (FMA) instructions

GPU Dialect

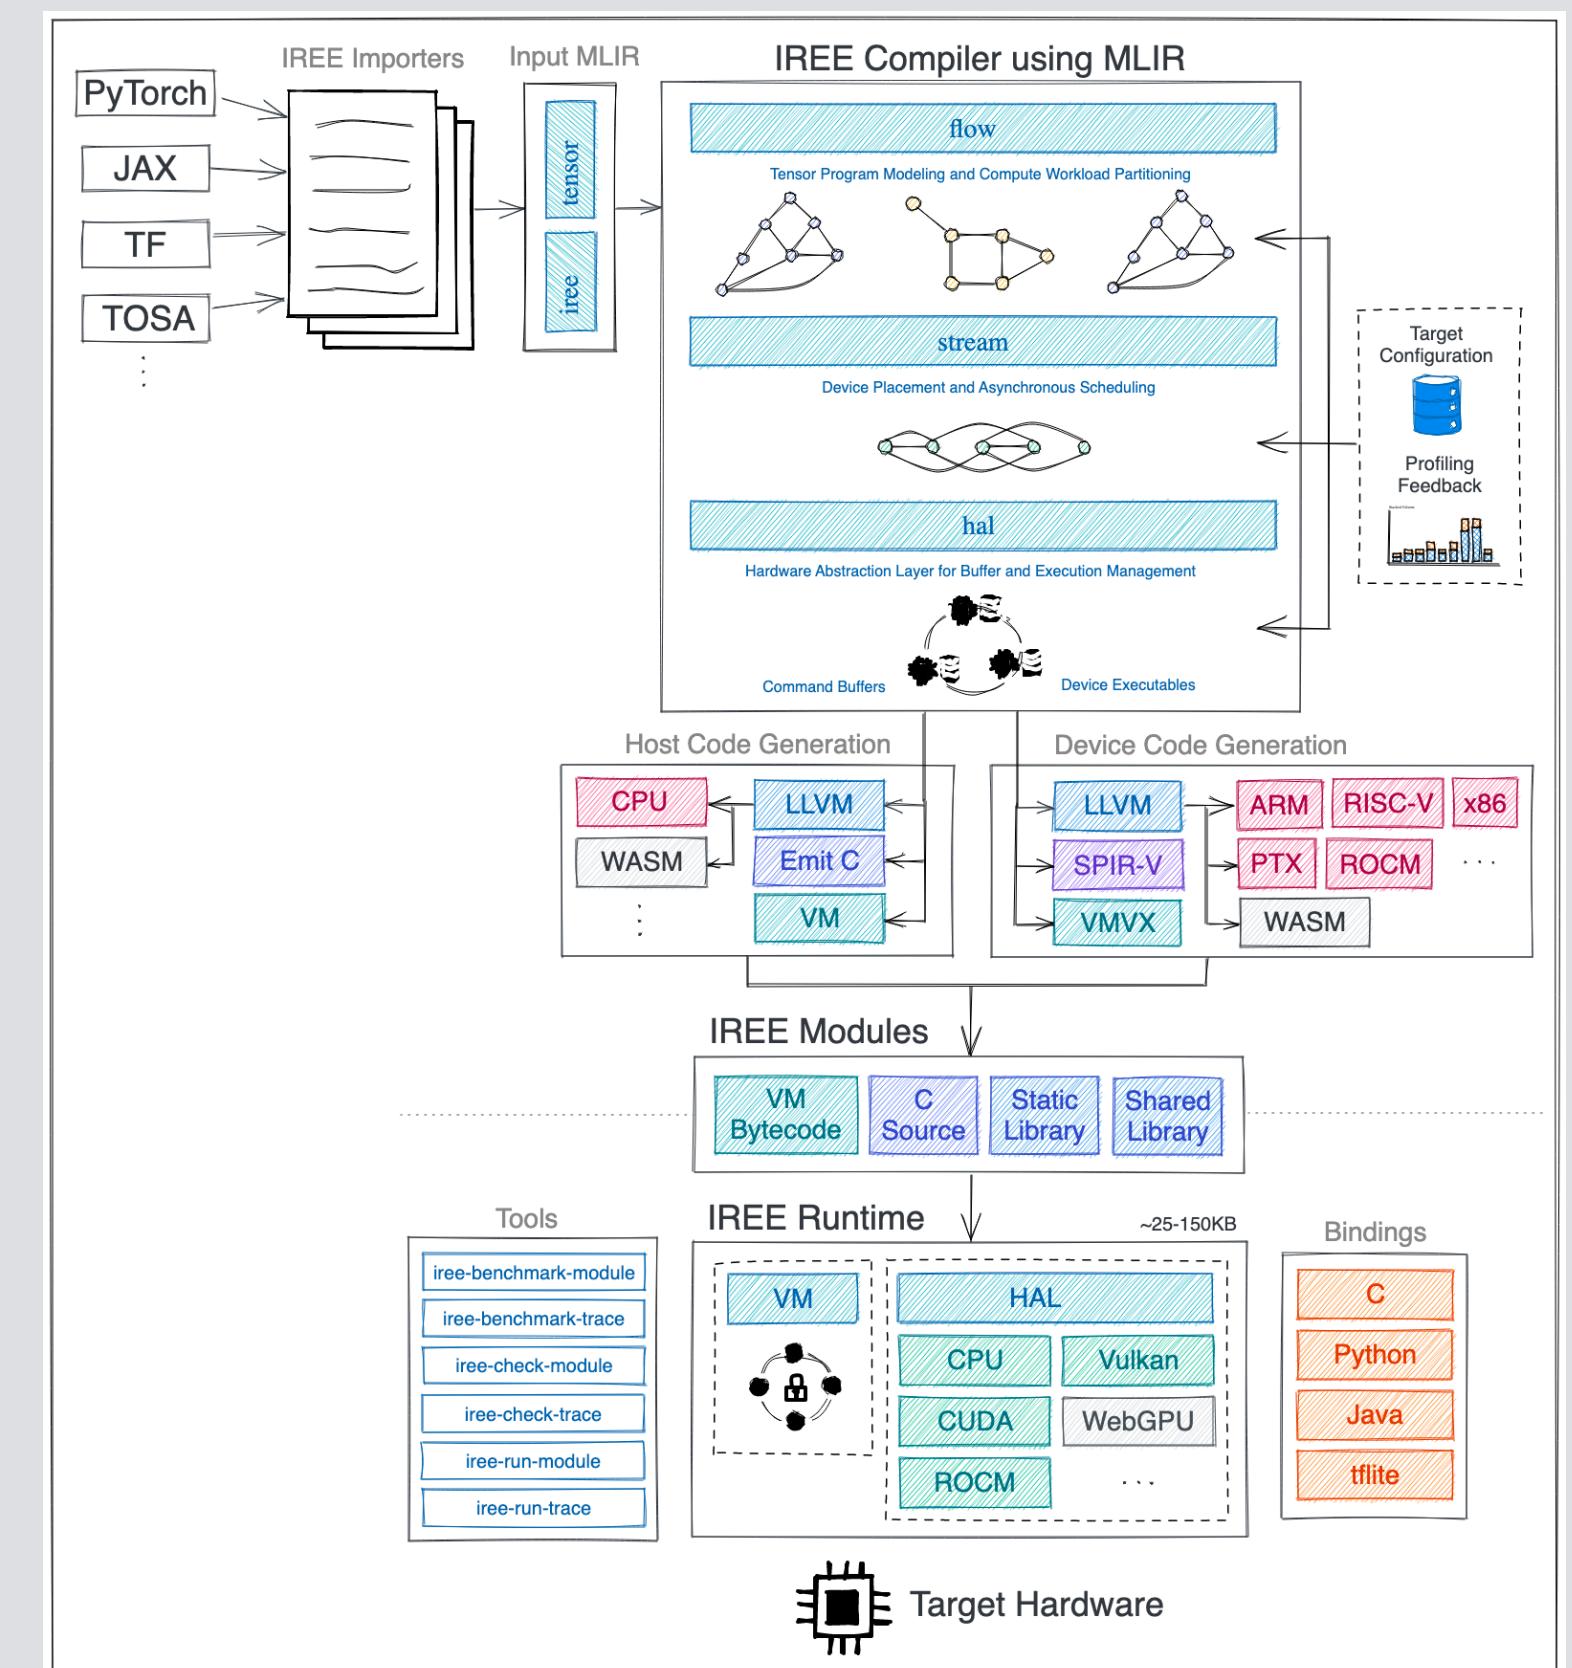
- Provides abstractions for retargetable GPU model
- Contains operations common to SIMT platforms
- Examples
 - Communication: `gpu.all_reduce` (reduction across a local workgroup)
 - Synchronization: `gpu.barrier`
 - Compute: `gpu.subgroup_mma_compute`
- Can be obtained by lowering from vector dialect
- Operates only on memrefs

Accelerator-Specific Dialects

- NVIDIA GPU specific dialects
 - NVVM dialect, NVGPU dialect
- AMD GPU specific dialects
 - ROCDL dialect
- ARM CPU specific dialects
 - ARM Neon dialect, ARM SVE dialect
- Intel CPU specific dialects
 - x86Vector dialect
- Each of these dialects exposes specific hardware functionality present in those devices
 - NVVM dialect has `nvvm.wmma.mma` for matrix multiplication using tensor cores, not present in other dialects

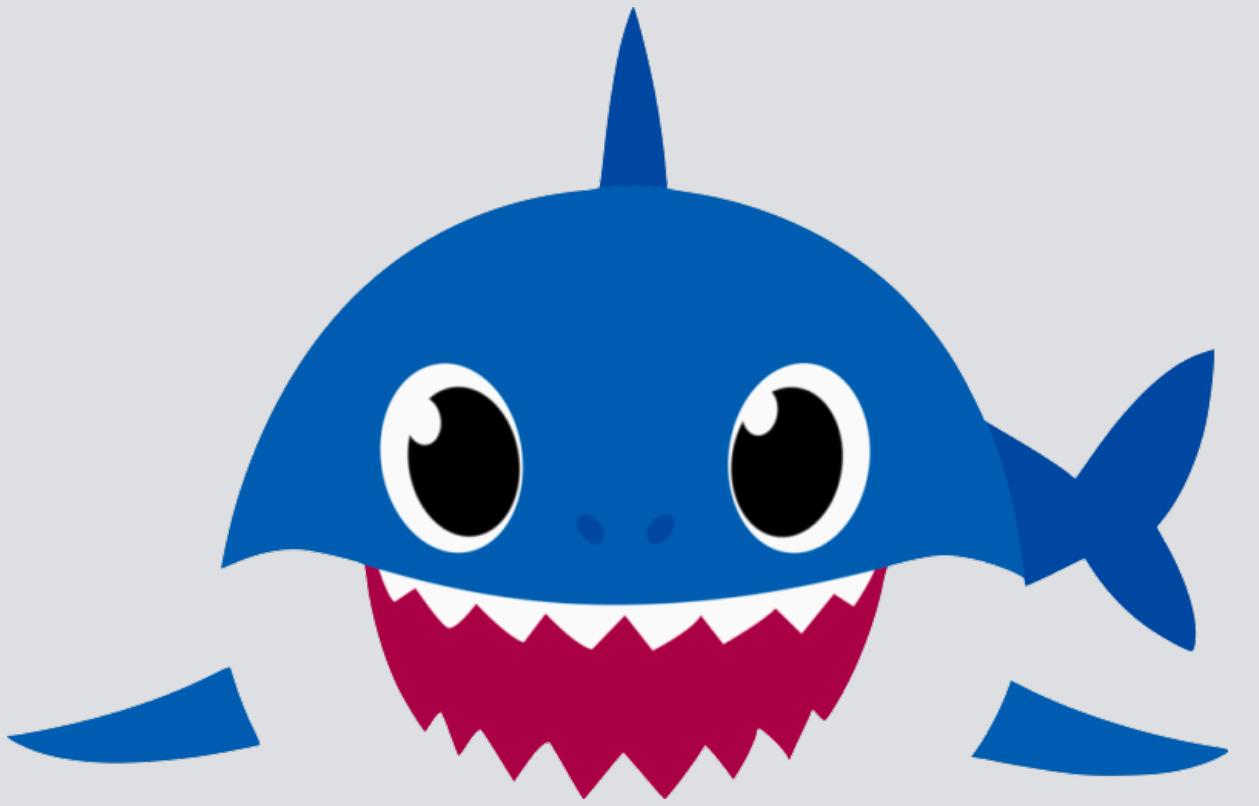
Code generation using IREE

- IREE is an open-source MLIR-based end-to-end compiler and runtime that lowers ML models for datacenter and edge workloads
- Supports X86, NVIDIA, AMD, RISC-V, Vulkan and ARM
- Supports Tensorflow, JAX, TFLite, PyTorch
- We will use IREE to demonstrate code generation using MLIR, specifically focusing on the compilation pipeline



Code generation using SHARK

- Builds on top of IREE pipeline
- Adds performance optimizations for CUDA
 - Caching allocator
 - Async memory prefetch
- Adds auto-tuning capabilities
- Adds backends for custom accelerators
- Contains a fully validated set of 100s of models
- Easy to deploy (integrated with triton server)



Perceptron Code Generation

- Perceptron is defined mathematically as shown below

$$z = \max(0, XW)$$

- In ML frameworks, it is a matrix multiplication followed by a ReLU non-linearity

- The code on the top right shows how to define such a network in **C++** and in **Tensorflow** on the bottom right

- By not having to start with a loop-based definition, we avoid any raising and perform optimizations on a ***higher level of abstraction*** than possible with other languages such as C++

- This results in reduced complexity during fusion, tiling and vectorization

```

for (int i = 0; i < 64; i++) {
    for (int j = 0; j < 1024; j++) {
        z[i][j] = 0.0;
        for (int k = 0; k < 1024; k++) {
            z[i][j] += x[i][k] * y[k][j];
        }
        z[i][j] = std::max(0.0, z[i][j]);
    }
}

```

```

BATCH_SIZE = 64
INPUT_SIZE = 1024
input_shapes = tf.TensorSpec(shape=[BATCH_SIZE, INPUT_SIZE],
                             dtype=tf.float32)

```

```

class MLP(tf.Module):
    def __init__(self) -> None:
        super(MLP, self).__init__()
        self.W = tf.Variable(tf.ones([INPUT_SIZE, INPUT_SIZE]))

```

```

@tf.function(input_signature=[input_shapes])
def predict(self, x: tf.Tensor) -> tf.Tensor:
    return tf.nn.relu(tf.matmul(x, self.W))

```

Lowering to Linalg on Tensors

- TF is lowered to MHLO and then to the linalg dialect
 - Both matrix multiplication and the ReLU are lowered to linalg.generic ops
 - Important to note that linalg.matmul is
- $$C_{i,j} = C_{i,j} + A_{i,k}B_{k,j}$$
- Hence, we need a linalg.fill before the linalg.matmul

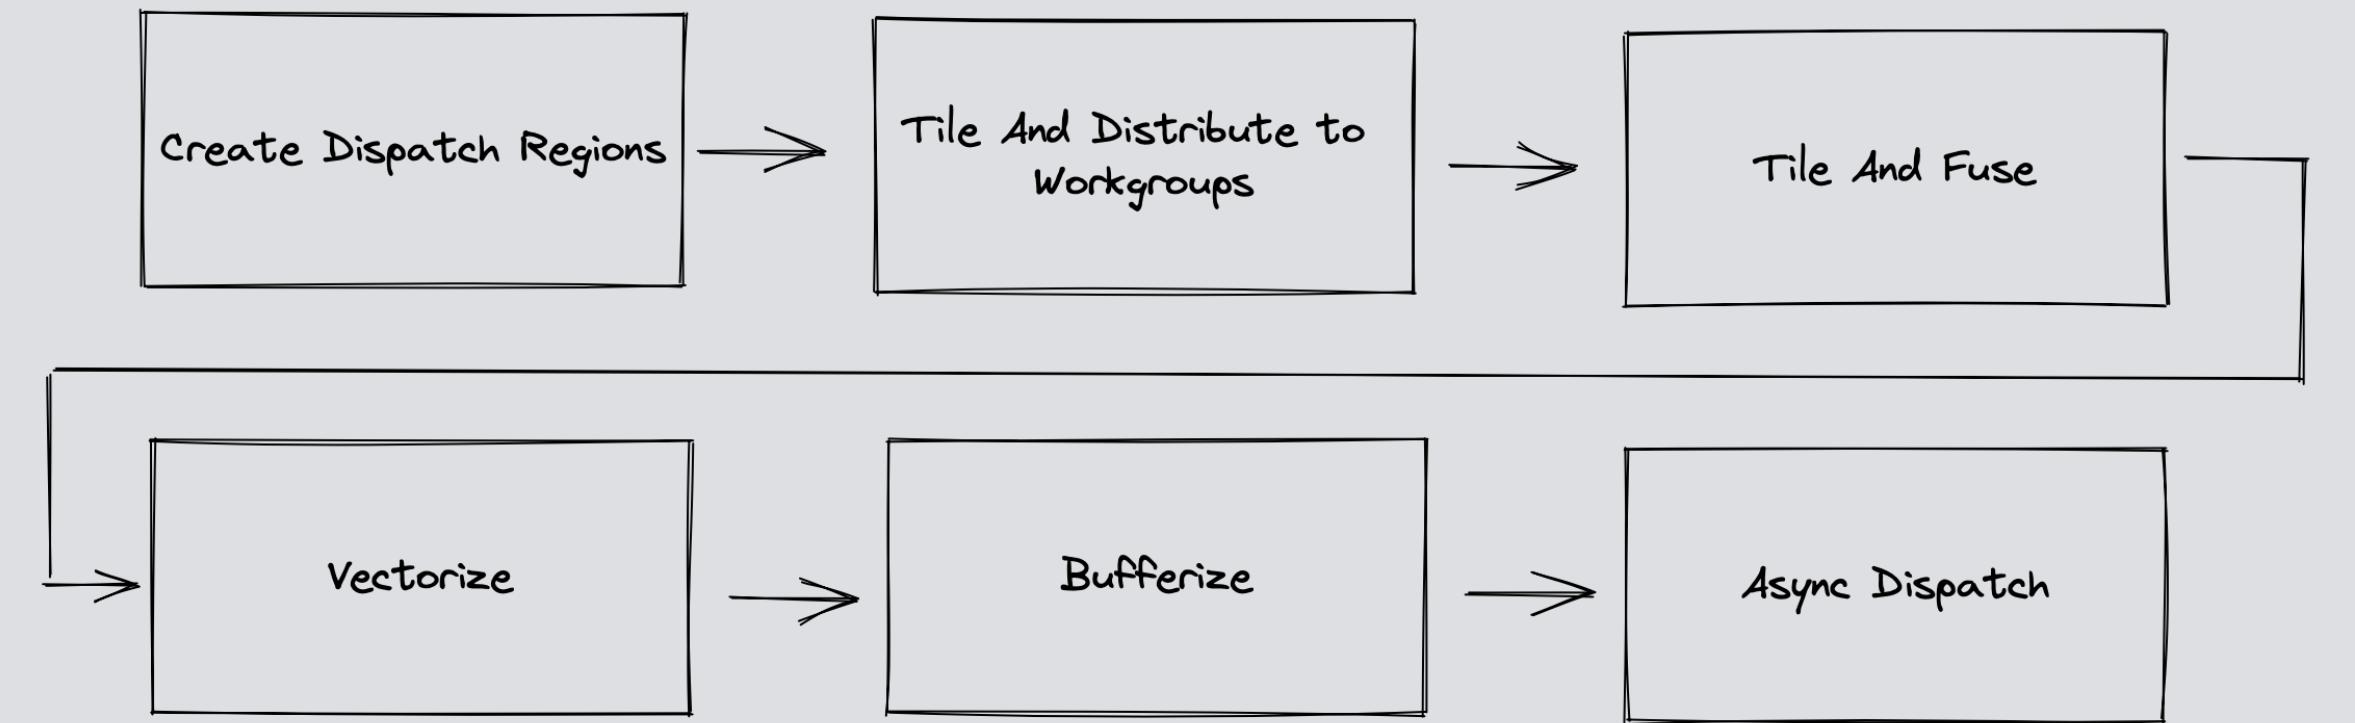
```

#map = affine_map<(d0, d1) -> (d0, d1)>
module {
    iree_input.global private @VV = dense<1.000000e+00> : tensor<1024x1024xf32>
    ...
    func.func private @___inference_predict_140(%arg0: tensor<64x1024xf32> {
        %cst = arith.constant 0.000000e+00 : f32
        ...
        %2 = linalg.init_tensor [64, 1024] : tensor<64x1024xf32>
        %3 = linalg.fill ins(%cst : f32) outs(%2 : tensor<64x1024xf32>) -> tensor<64x1024xf32>
        %4 = linalg.matmul ins(%arg0, %1 : tensor<64x1024xf32>, tensor<1024x1024xf32>)
                                         outs(%3 : tensor<64x1024xf32>) -> tensor<64x1024xf32>
        %5 = linalg.init_tensor [64, 1024] : tensor<64x1024xf32>
        %6 = linalg.generic {indexing_maps = [#map, #map], iterator_types = ["parallel", "parallel"]}
              ins(%4 : tensor<64x1024xf32>) outs(%5 : tensor<64x1024xf32>) {
                  ^bb0(%arg1: f32, %arg2: f32):
                      %7 = arith.maxf %arg1, %cst : f32
                      linalg.yield %7 : f32
              } -> tensor<64x1024xf32>
        return %6 : tensor<64x1024xf32>
    }
}

```

Overview of CPU code generation pipeline

- Split computation graph into smaller sub-graphs (dispatch regions)
- Tiling and vectorization key pieces of CPU pipeline
- CPU pipeline performs bufferization late, preferring in-place bufferization while taking care to avoid RaW conflicts
- Finally operations are lowered to make mapping to hardware instructions efficient (such as lowering matrix multiplication to outer products to easily map to FMA instructions)
- Most of this pipeline shared with GPU



Dispatch Region Formation

- Dispatch region contains computation that has to be executed on device in an atomic fashion
- Large neural networks are partitioned into finite number of dispatch regions
- Each dispatch region contains a root op (any linalg named op or generic op with reduction iterator type)
- Root ops are then fused with consumers if all uses of producer are dominated by it
- Also does elementwise fusion to fuse linalg.generic ops

```

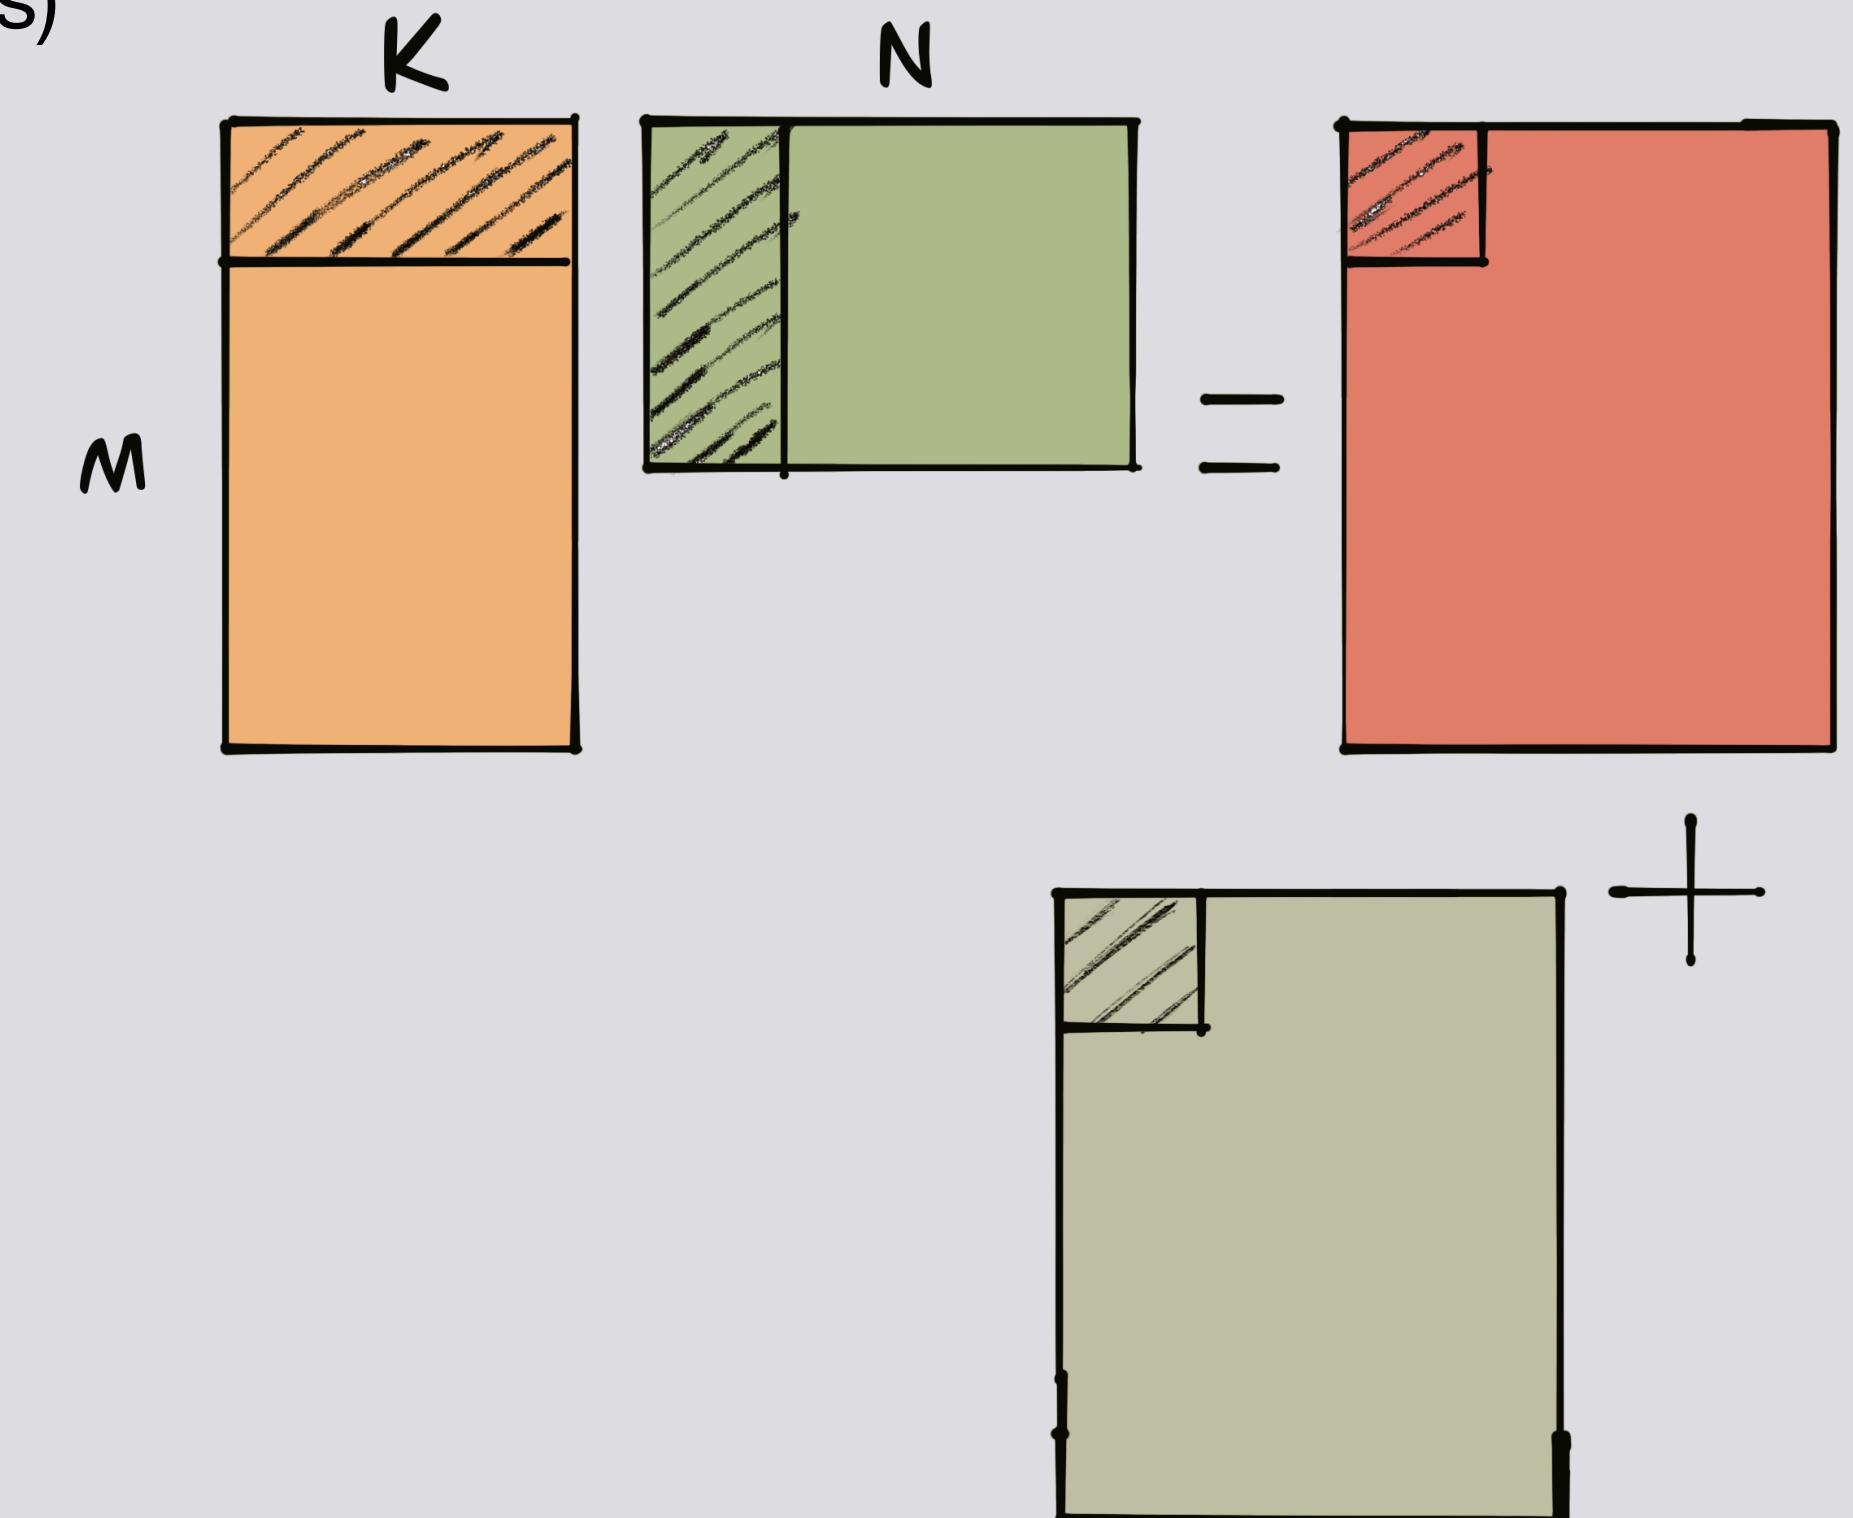
%1 = flow.dispatch.workgroups[%c64, %c1024](%0) : (tensor<64x1024xf32>) -> tensor<64x1024xf32>
(%arg1: !flow.dispatch.tensor<readonly:64x1024xf32>, %arg2:
 !flow.dispatch.tensor<writeonly:64x1024xf32>) {
%3 = flow.dispatch.tensor.load %arg1, offsets = [0, 0], sizes = [64, 1024], strides = [1, 1] :
    !flow.dispatch.tensor<readonly:64x1024xf32> -> tensor<64x1024xf32>
%4 = linalg.init_tensor [64, 1024] : tensor<64x1024xf32>
%5 = linalg.fill ins(%cst : f32) outs(%4 : tensor<64x1024xf32>) -> tensor<64x1024xf32>
%6 = linalg.matmul ins(%3, %cst_0 : tensor<64x1024xf32>, tensor<1024x1024xf32>)
    outs(%5 : tensor<64x1024xf32>) -> tensor<64x1024xf32>
%7 = linalg.generic {indexing_maps = [affine_map<(d0, d1) -> (d0, d1)>,
    affine_map<(d0, d1) -> (d0, d1)>],
    iterator_types = ["parallel", "parallel"]}
ins(%6 : tensor<64x1024xf32>) outs(%4 : tensor<64x1024xf32>) {
^bb0(%arg3: f32, %arg4: f32):
    %8 = arith.maxf %arg3, %cst : f32
    linalg.yield %8 : f32
} -> tensor<64x1024xf32>
flow.dispatch.tensor.store %7, %arg2, offsets = [0, 0], sizes = [64, 1024], strides = [1, 1] :
    tensor<64x1024xf32> -> !flow.dispatch.tensor<writeonly:64x1024xf32>
flow.return
} count(%arg1: index, %arg2: index) -> (index, index, index) {
    %x, %y, %z = flow.dispatch.default_workgroup_count %arg1, %arg2
    flow.return %x, %y, %z : index, index, index
}

```

Body of Dispatch Region

Tile and Distribute to Workgroups

- Work partitioned along a 3-D grid of virtual processors (workgroups) that can be mapped to multi-core CPUs or GPUs
- Uses a block cyclic distribution to distribute the tiles
- Each dispatch region performs a tile of the computation
- The compute done is determined by its rank (`workgroup_id_x`, `workgroup_id_y`) and number of processors
- Only parallel dimensions are tiled

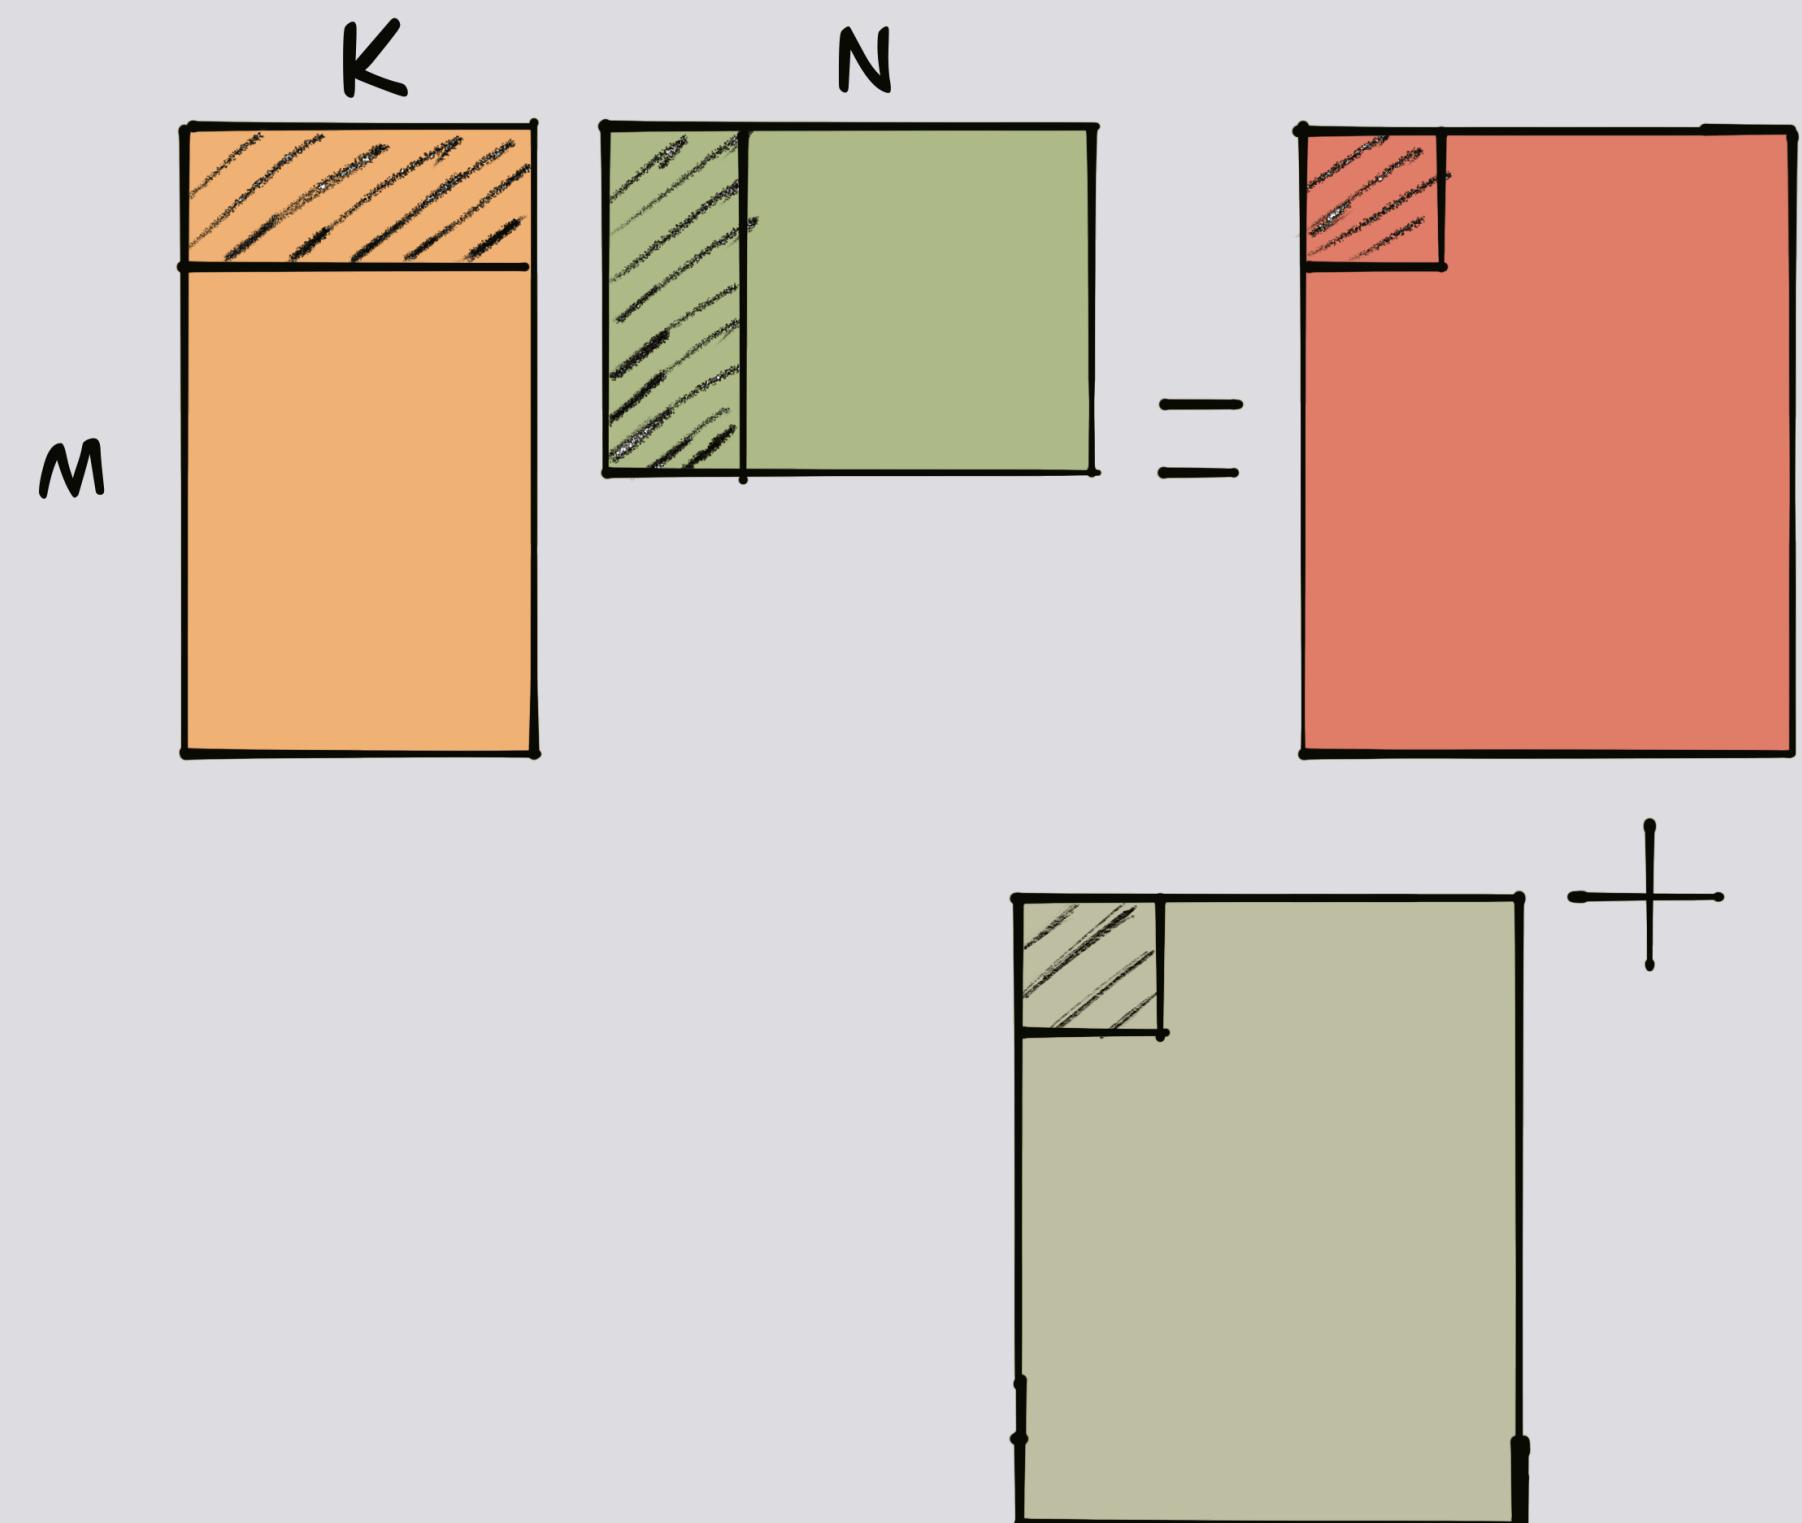


Tile and Distribute to Workgroups

```

func.func @predict_dispatch_00 {
    ...
    %workgroup_id_x = hal.interface.workgroup.id[0] : index
    %workgroup_count_x = hal.interface.workgroup.count[0] : index
    %workgroup_id_y = hal.interface.workgroup.id[1] : index
    %workgroup_count_y = hal.interface.workgroup.count[1] : index
    %2 = affine.apply affine_map<0[s0] -> (s0 * 32)>0[%workgroup_id_y]
    %3 = affine.apply affine_map<0[s0] -> (s0 * 32)>0[%workgroup_count_y]
    scf.for %arg0 = %2 to %c64 step %3 {
        %4 = affine.apply affine_map<0[s0] -> (s0 * 32)>0[%workgroup_id_x]
        %5 = affine.apply affine_map<0[s0] -> (s0 * 32)>0[%workgroup_count_x]
        scf.for %arg1 = %4 to %c1024 step %5 {
            ...
            %8 = linalg.fill ins(%cst_0 : f32) outs(%6 : tensor<32x32xf32>) -> tensor<32x32xf32>
            %9 = linalg.matmul ins(%7, %cst : tensor<32x1024xf32>, tensor<1024x32xf32>
                outs(%8 : tensor<32x32xf32>) -> tensor<32x32xf32>
            %10 = linalg.generic {indexing_maps = [affine_map<(d0, d1) -> (d0, d1)>,
                iterator_types = ["parallel", "parallel"]]} outs(%9 : tensor<32x32xf32>) {
                ^bb0(%arg2: f32):
                    %11 = arith.maxf %arg2, %cst_0 : f32
                    linalg.yield %11 : f32
            } -> tensor<32x32xf32>
            ...
        }
    }
}

```



Tile and Fuse

- Additional tiling along parallel dimensions
- Introduces subset operations (`tensor.extract_slice`, `tensor.insert_slice`) to access the tiled data
- Depending on chosen tile sizes, there may not exists a single static tensor type valid for every iteration
- The sub-tensor may be relaxed to a dynamic tensor
- Subsequent canonicalizations can be used to refine any shapes that are determined to be static
- Can use padding to handle these scenarios by introducing `tensor.pad` operations with appropriate sizes

```

func.func @predict_dispatch_00 {
  scf.for %arg0 = %2 to %c64 step %3 {
    scf.for %arg1 = %4 to %c1024 step %5 {
      %8 = scf.for %arg2 = %c0 to %c32 step %c8 iter_args(%arg3 = %7) -> (tensor<32x128xf32>) {
        %9 = tensor.extract_slice %6[%arg2, 0] [8, 1024] [1, 1] : tensor<32x1024xf32> to
          tensor<8x1024xf32>
        %10 = scf.for %arg4 = %c0 to %c128 step %c32 iter_args(%arg5 = %arg3) -> (tensor<32x128xf32>) {
          %11 = tensor.extract_slice %arg5[%arg2, %arg4] [8, 32] [1, 1] : tensor<32x128xf32>
            to tensor<8x32xf32>
          %12 = linalg.fill ins(%cst_0 : f32) outs(%11 : tensor<8x32xf32>) -> tensor<8x32xf32>
          %13 = linalg.matmul ins(%9, %cst : tensor<8x1024xf32>, tensor<1024x32xf32>)
            outs(%12 : tensor<8x32xf32>) -> tensor<8x32xf32>
          %14 = linalg.generic outs(%13 : tensor<8x32xf32>) {
            ^bb0(%arg6: f32):
              %16 = arith.maxf %arg6, %cst_0 : f32
              linalg.yield %16 : f32
            } -> tensor<8x32xf32>
          %15 = tensor.insert_slice %14 into %arg5[%arg2, %arg4] [8, 32] [1, 1] : tensor<8x32xf32> into
            tensor<32x128xf32>
          scf.yield %15 : tensor<32x128xf32>
        }
        scf.yield %10 : tensor<32x128xf32>
      }
    }
  }
}

```

Single Tiling Expert

- Multiple different ways of tiling the matmul
 - DoubleTilingExpert, TripleTilingExpert, DoubleTilingPadExpert etc.
- Using SingleTilingExpert get a single tiling of reduction dimension
- Changes K dimension from 1024 to 16

```

func.func @predict_dispatch_00 {
  scf.for %arg0 = %2 to %c64 step %3 {
    scf.for %arg1 = %4 to %c1024 step %5 {
      %8 = scf.for %arg2 = %c0 to %c32 step %c8 iter_args(%arg3 = %7) -> (tensor<32x128xf32>) {
        %9 = scf.for %arg4 = %c0 to %c128 step %c32 iter_args(%arg5 = %arg3) -> (tensor<32x128xf32>) {
          %10 = tensor.extract_slice %arg5[%arg2, %arg4] [8, 32] [1, 1] : tensor<32x128xf32>
              to tensor<8x32xf32>

          %11 = linalg.fill ins(%cst_0 : f32) outs(%10 : tensor<8x32xf32>) -> tensor<8x32xf32>
          %12 = scf.for %arg6 = %c0 to %c1024 step %c16 iter_args(%arg7 = %11) -> (tensor<8x32xf32>) {
            %15 = tensor.extract_slice %6[%arg2, %arg6] [8, 16] [1, 1] : tensor<32x1024xf32> to tensor<8x16xf32>
            %16 = linalg.matmul ins(%15, %cst : tensor<8x16xf32>, tensor<16x32xf32>
                                outs(%arg7 : tensor<8x32xf32>) -> tensor<8x32xf32>
            scf.yield %16 : tensor<8x32xf32>
          }

          %13 = linalg.generic ...
          %14 = tensor.insert_slice %13 into %arg5[%arg2, %arg4] [8, 32] [1, 1] : tensor<8x32xf32>
              into tensor<32x128xf32>
          scf.yield %14 : tensor<32x128xf32>
        }

        scf.yield %9 : tensor<32x128xf32>
      }
    }
  }
}

```

Vectorize

- Emits `vector.transfer_read/write` operations for each operand
- For elementwise operations, rewrite as pointwise vector variant
- For reductions, rewrite as `vector.contract` or multi-reduction
- Broadcasting lower dimensional operands is done by `vector.broadcast`
- Permutations are handled by `vector.transpose`
- Lowers n-D vectors to 1-D vectors supported by LLVM

```

func.func @predict_dispatch_00 {
%6 = vector.transfer_read %cst_0[%c0, %c0], %cst_1 {in_bounds = [true, true]} :
  tensor<16x32xf32>, vector<16x32xf32>
scf.for %arg0 = %2 to %c64 step %3 {
  scf.for %arg1 = %4 to %c1024 step %5 {
    %9 = scf.for %arg2 = %c0 to %c32 step %c8 iter_args(%arg3 = %8) -> (tensor<32x128xf32>) {
      %10 = scf.for %arg4 = %c0 to %c128 step %c32 iter_args(%arg5 = %arg3) -> (tensor<32x128xf32>) [
        %11 = scf.for %arg6 = %c0 to %c1024 step %c16 iter_args(%arg7 = %cst) -> (vector<8x32xf32>) [
          %14 = vector.transfer_read %7[%arg2, %arg6], %cst_1 {in_bounds = [true, true]} :
            tensor<32x1024xf32>, vector<8x16xf32>
          %15 = vector.contract {indexing_maps = [affine_map<(d0, d1, d2) -> (d0, d2)>,
                                              affine_map<(d0, d1, d2) -> (d2, d1)>,
                                              affine_map<(d0, d1, d2) -> (d0, d1)>],
                                 iterator_types = ["parallel", "parallel", "reduction"],
                                 kind = #vector.kind<add>}
            %14, %6, %arg7 : vector<8x16xf32>, vector<16x32xf32> into vector<8x32xf32>
          scf.yield %15 : vector<8x32xf32>
        }
      %12 = arith.maxf %11, %cst : vector<8x32xf32>
      %13 = vector.transfer_write %12, %arg5[%arg2, %arg4] {in_bounds = [true, true]} :
        vector<8x32xf32>, tensor<32x128xf32>
      scf.yield %13 : tensor<32x128xf32>
    }
    scf.yield %10 : tensor<32x128xf32>
  }
}

```

Bufferize

- Allocate and copy as little memory as possible
- Always prefer re-using buffers in place
- Use destination-passing style as a heuristic for in-place bufferization
- Tie output tensor to results tensor to act as bufferization constraint
- Performs a future in-place bufferization analysis of the operands and checks if a RaW conflict is detected
- If not, then performs in-place bufferization

```

func.func @predict_dispatch_00 {
    %7 = vector.transfer_read %0[%c0, %c0], %cst {in_bounds = [true, true]} : memref<16x32xf32>, vector<16x32xf32>

    scf.for %arg0 = %3 to %c64 step %4 {
        %8 = memref.subview %1[%arg0, 0] [32, 1024] [1, 1] : memref<64x1024xf32> to memref<32x1024xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>
        scf.for %arg1 = %5 to %c1024 step %6 {
            %9 = memref.subview %2[%arg0, %arg1] [32, 128] [1, 1] : memref<64x1024xf32> to memref<32x128xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>
                scf.for %arg2 = %c0 to %c32 step %c8 {
                    scf.for %arg3 = %c0 to %c128 step %c32 {
                        %10 = scf.for %arg4 = %c0 to %c1024 step %c16 iter_args(%arg5 = %cst_0) -> (vector<8x32xf32>)
                            %12 = vector.transfer_read %8[%arg2, %arg4], %cst {in_bounds = [true, true]} : memref<32x1024xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>, vector<8x16xf32>
                            %13 = vector.contract {indexing_maps = [affine_map<(d0, d1, d2) -> (d0, d2)>, affine_map<(d0, d1, d2) -> (d2, d1)>, affine_map<(d0, d1, d2) -> (d0, d1)>], iterator_types = ["parallel", "parallel", "reduction"], kind = #vector.kind<add>}
                                %12, %7, %arg5 : vector<8x16xf32>, vector<16x32xf32> into vector<8x32xf32>
                            scf.yield %13 : vector<8x32xf32>
                        }
                        %11 = arith.maxf %10, %cst_0 : vector<8x32xf32>
                        vector.transfer_write %11, %9[%arg2, %arg3] {in_bounds = [true, true]} : vector<8x32xf32>, memref<32x128xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>
                    }
                }
            }
        }
    }
}

```

Lowering closer to hardware

- Apply vector unrolling
 - Breaks down vector sizes to sizes well supported by target
 - Pre-emptively handles non power of 2 sizes to avoid suboptimal code generation
- `vector.contract` is lowered to outer products to enable mapping to SIMD FMA instructions
- Further lowered to LLVM dialect and translated to LLVM IR

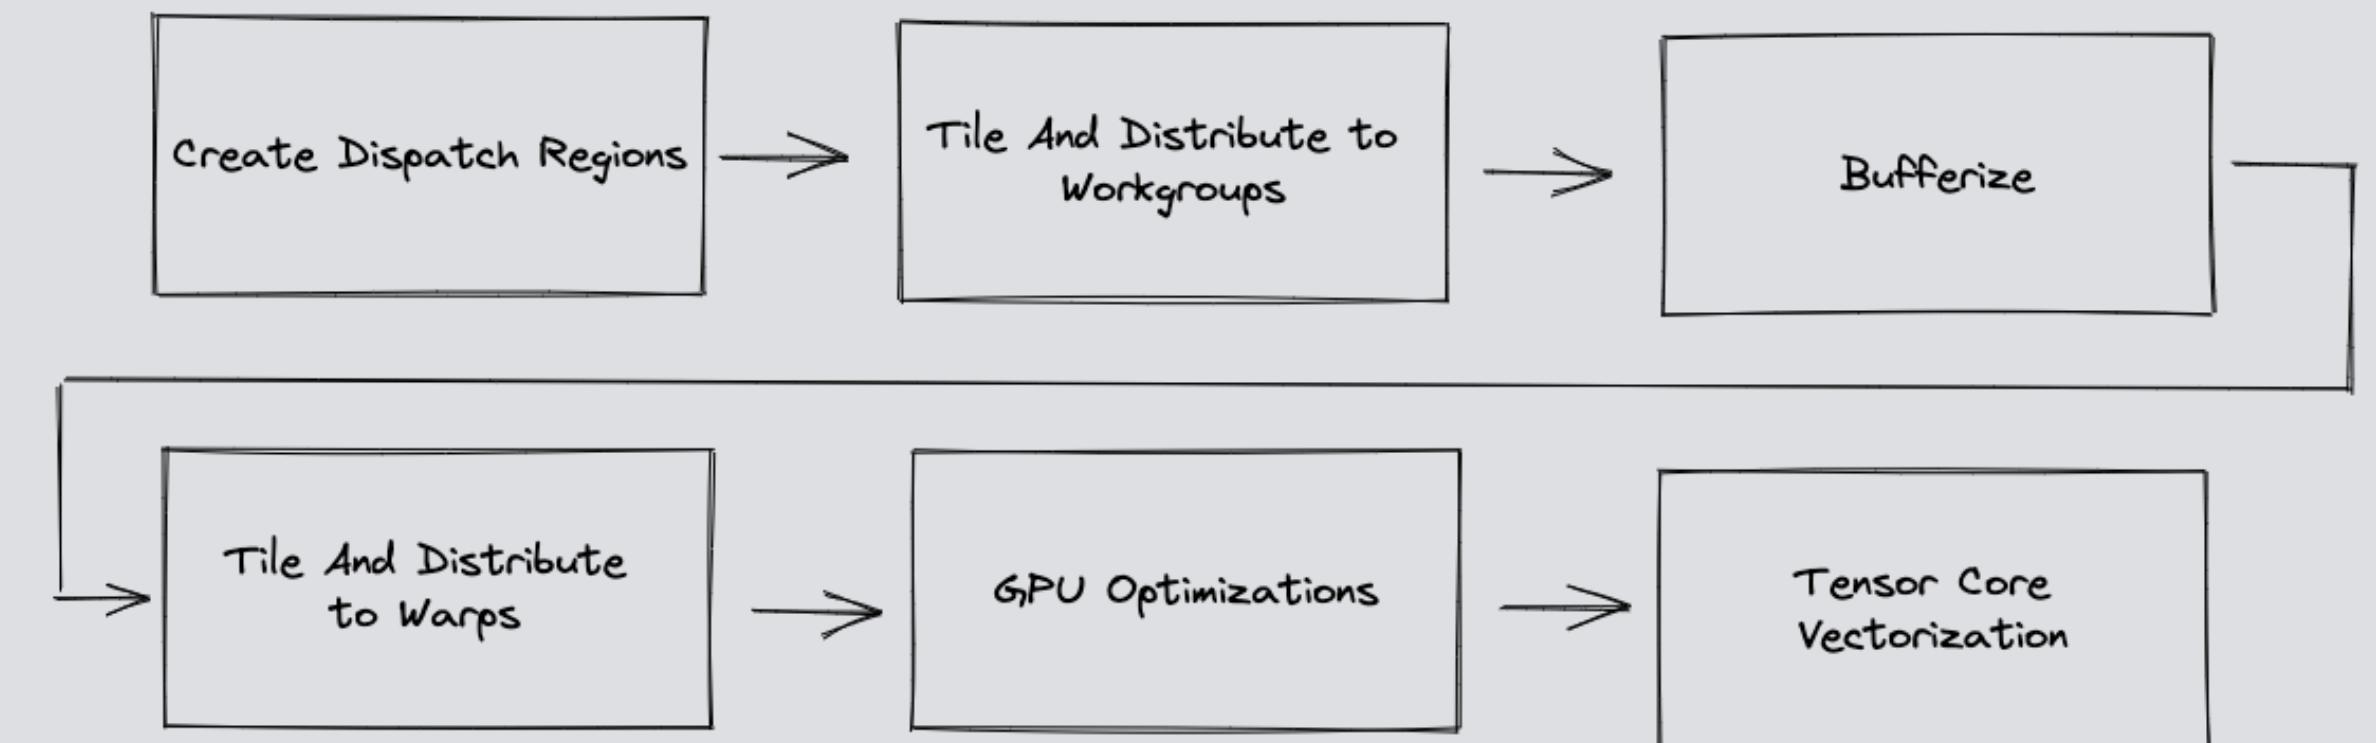
```

func.func @predict_dispatch_00 {
  ...
  %8 = vector.extract %5[0] : vector<16x32xf32>
  %9 = vector.extract %5[1] : vector<16x32xf32>
  ...
  %21 = vector.extract %5[13] : vector<16x32xf32>
  %22 = vector.extract %5[14] : vector<16x32xf32>
  scf.for %arg0 = %c0 to %c32 step %c8 {
    scf.for %arg1 = %c0 to %c128 step %c32 {
      %24 = scf.for %arg2 = %c0 to %c1024 step %c16 iter_args(%arg3 = %cst_0) -> (vector<8x32xf32>) {
        %26 = vector.transfer_read %6[%arg0, %arg2], %cst {in_bounds = [true, true]} :
               memref<32x1024xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>, vector<8x16xf32>
        %27 = vector.transpose %26, [1, 0] : vector<8x16xf32> to vector<16x8xf32>
        %28 = vector.extract %27[0] : vector<16x8xf32>
        %29 = vector.outerproduct %28, %8, %arg3 {kind = #vector.kind<add>} : vector<8xf32>, vector<32xf32>
        ...
        %58 = vector.extract %27[15] : vector<16x8xf32>
        %59 = vector.outerproduct %58, %23, %57 {kind = #vector.kind<add>} : vector<8xf32>, vector<32xf32>
        scf.yield %59 : vector<8x32xf32>
    }
    %25 = arith.maxf %24, %cst_0 : vector<8x32xf32>
    vector.transfer_write %25, %7[%arg0, %arg1] {in_bounds = [true, true]} : vector<8x32xf32>,
               memref<32x128xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>
  }
}

```

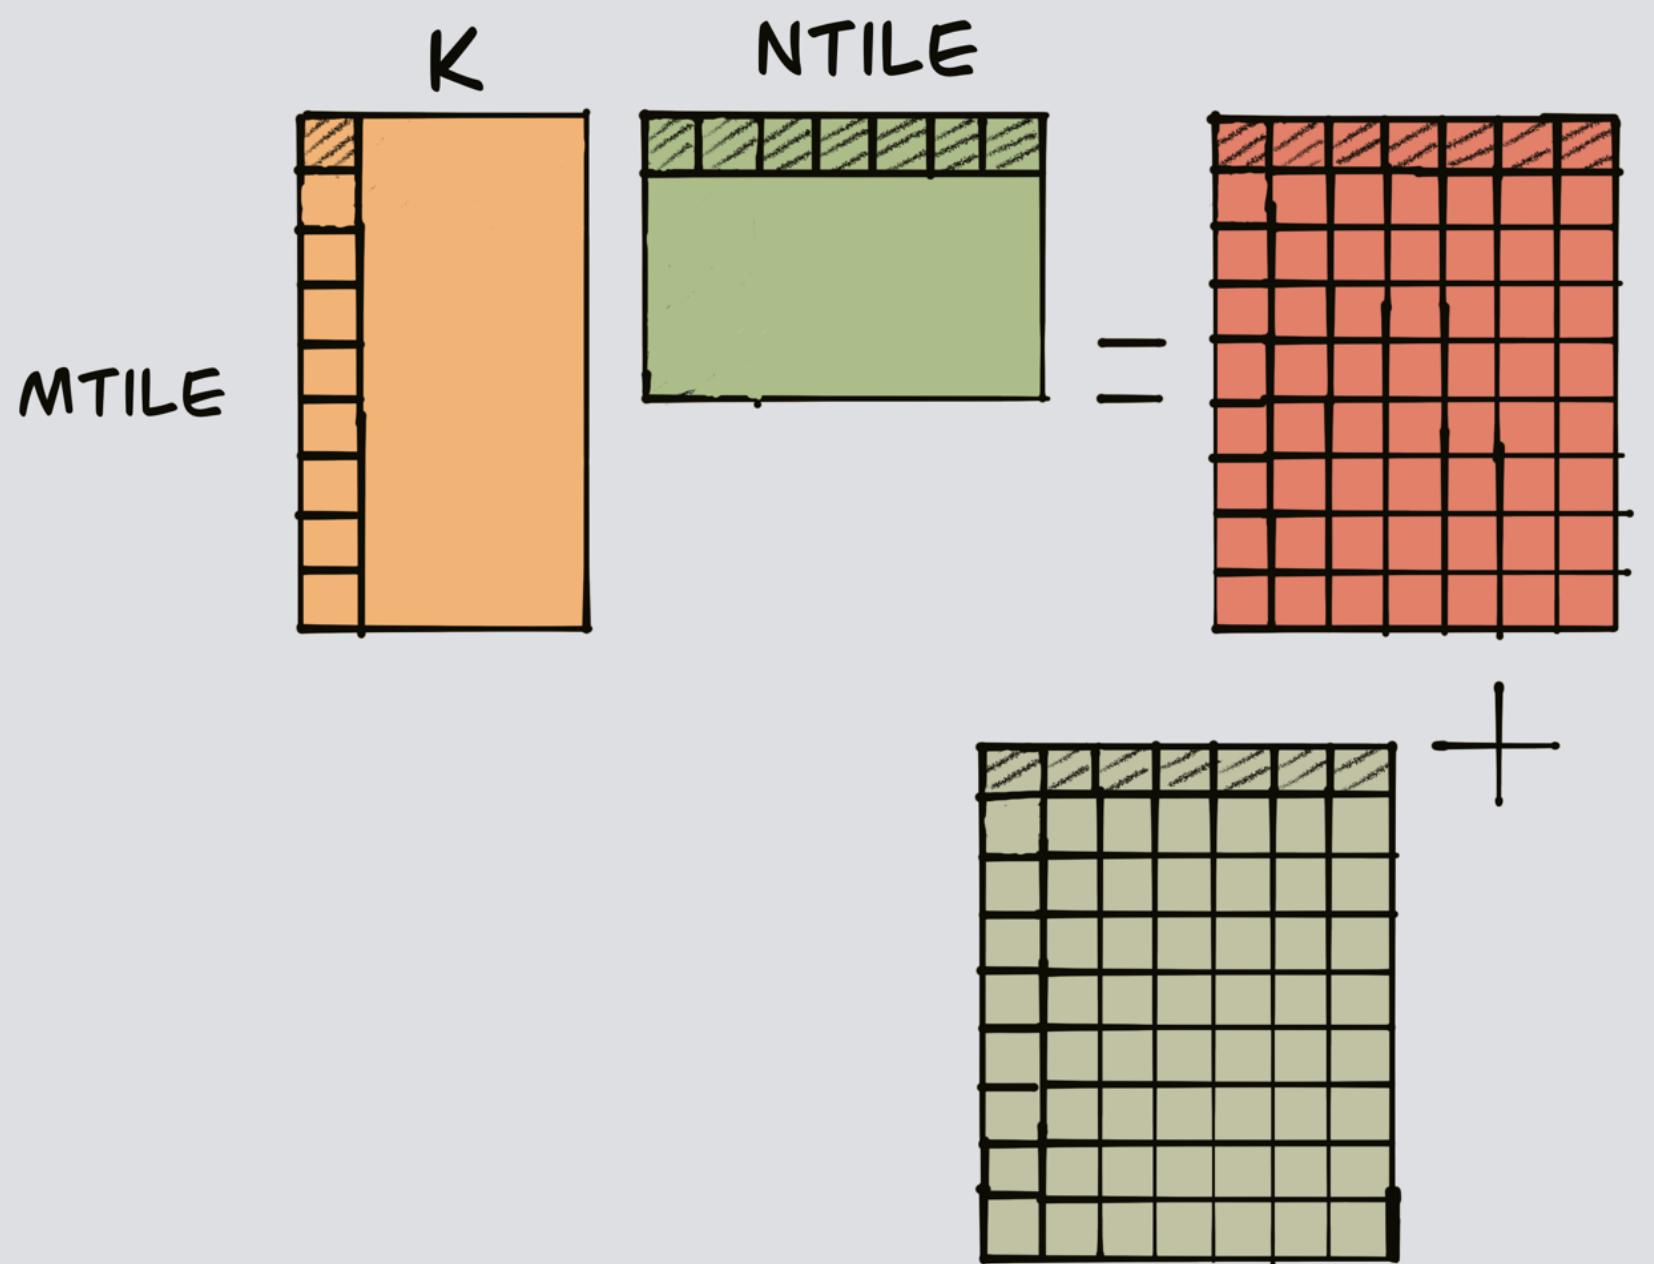
Overview of GPU code generation pipeline

- Tiling and vectorization key pieces of GPU pipeline
- GPU pipeline bufferizes early and focuses on optimizing shared memory copies and reducing bank conflicts
- Plan to move bufferization after vectorization
- GPU vectorization focuses on using tensor cores efficiently and emitting appropriate nvvm intrinsics



Tile and Distribute to Warps

- Tile and distribute to warps
 - After bufferization, tile the reduction dimension
 - Copy subviews of input memrefs to shared memory (workgroup memory) on the GPU prior to computation
 - Start introducing SIMD characteristics by introducing `memref.copy` (mapping to shared memory done later downstream)
 - Insert barriers after copying to workgroup memory
 - Do an additional level of tiling to distribute to a warps

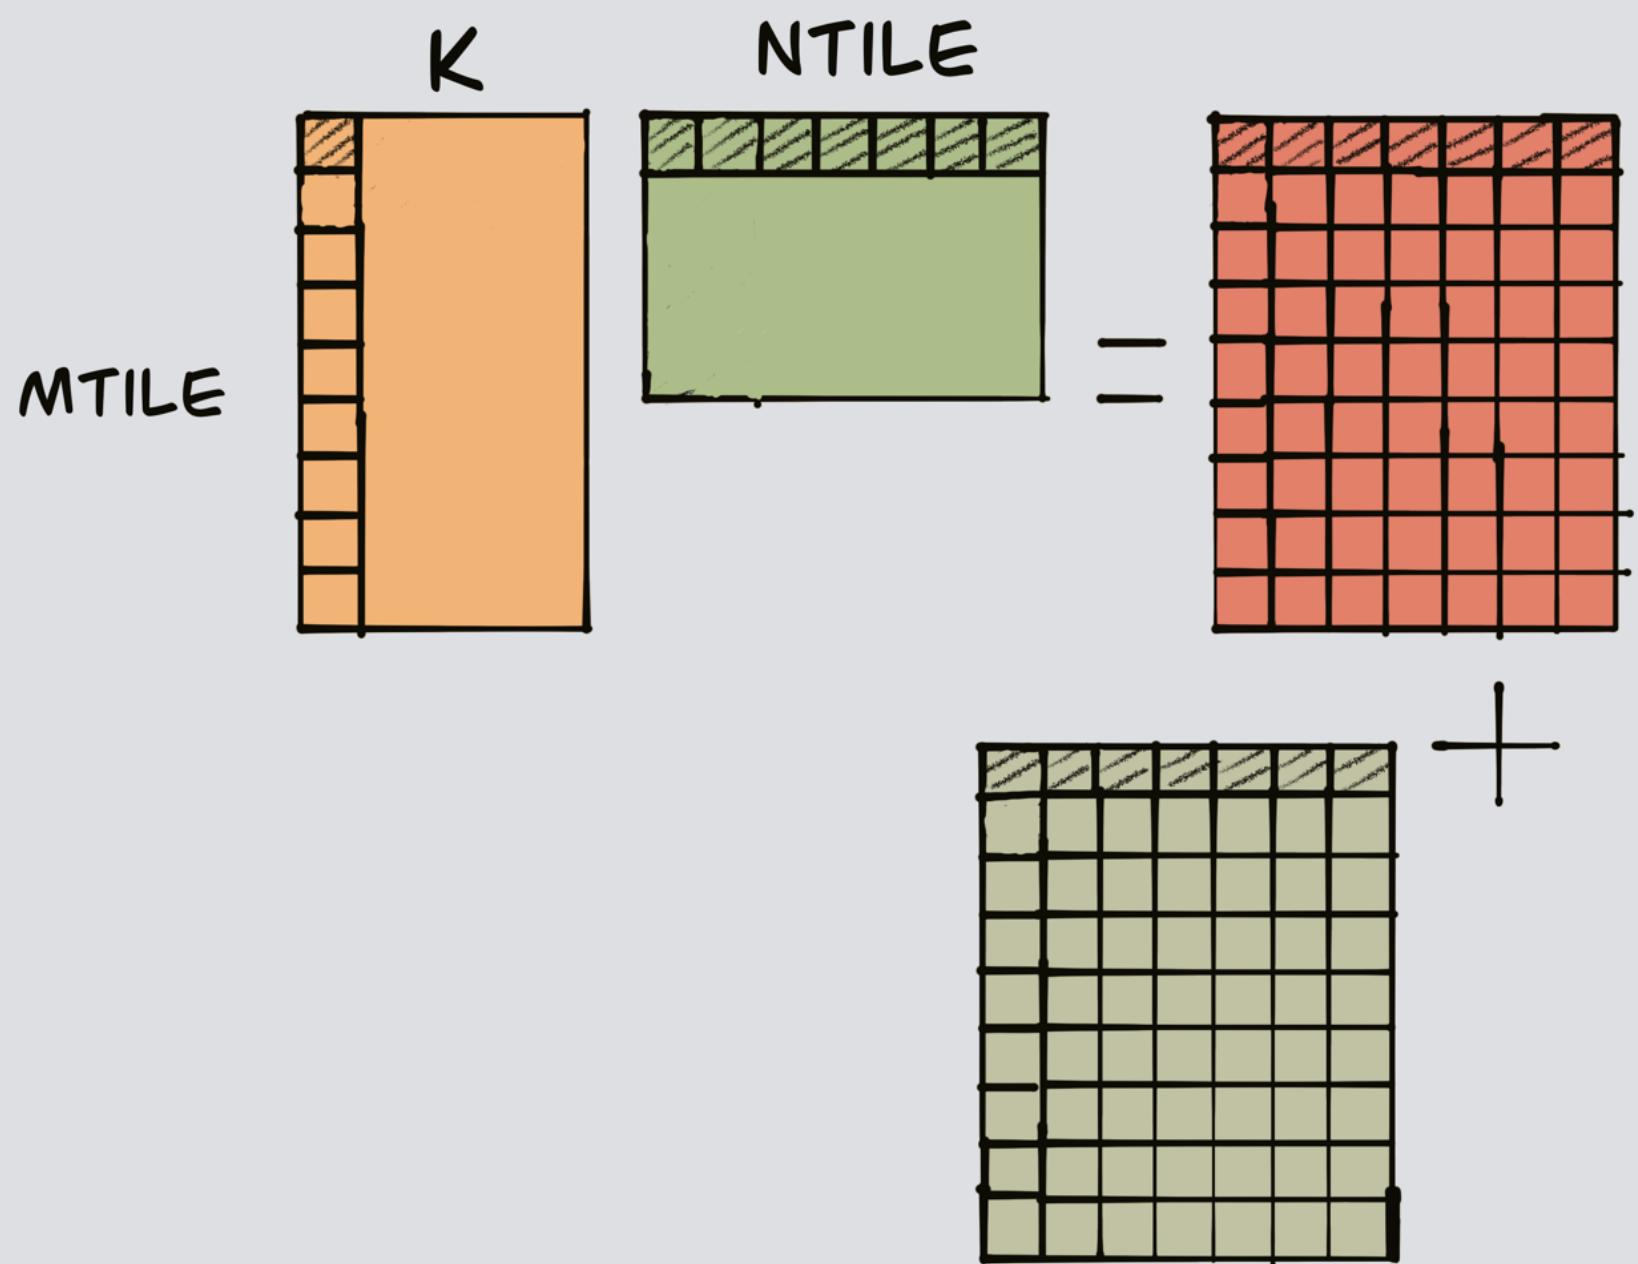


Tile and Distribute to Warps

```

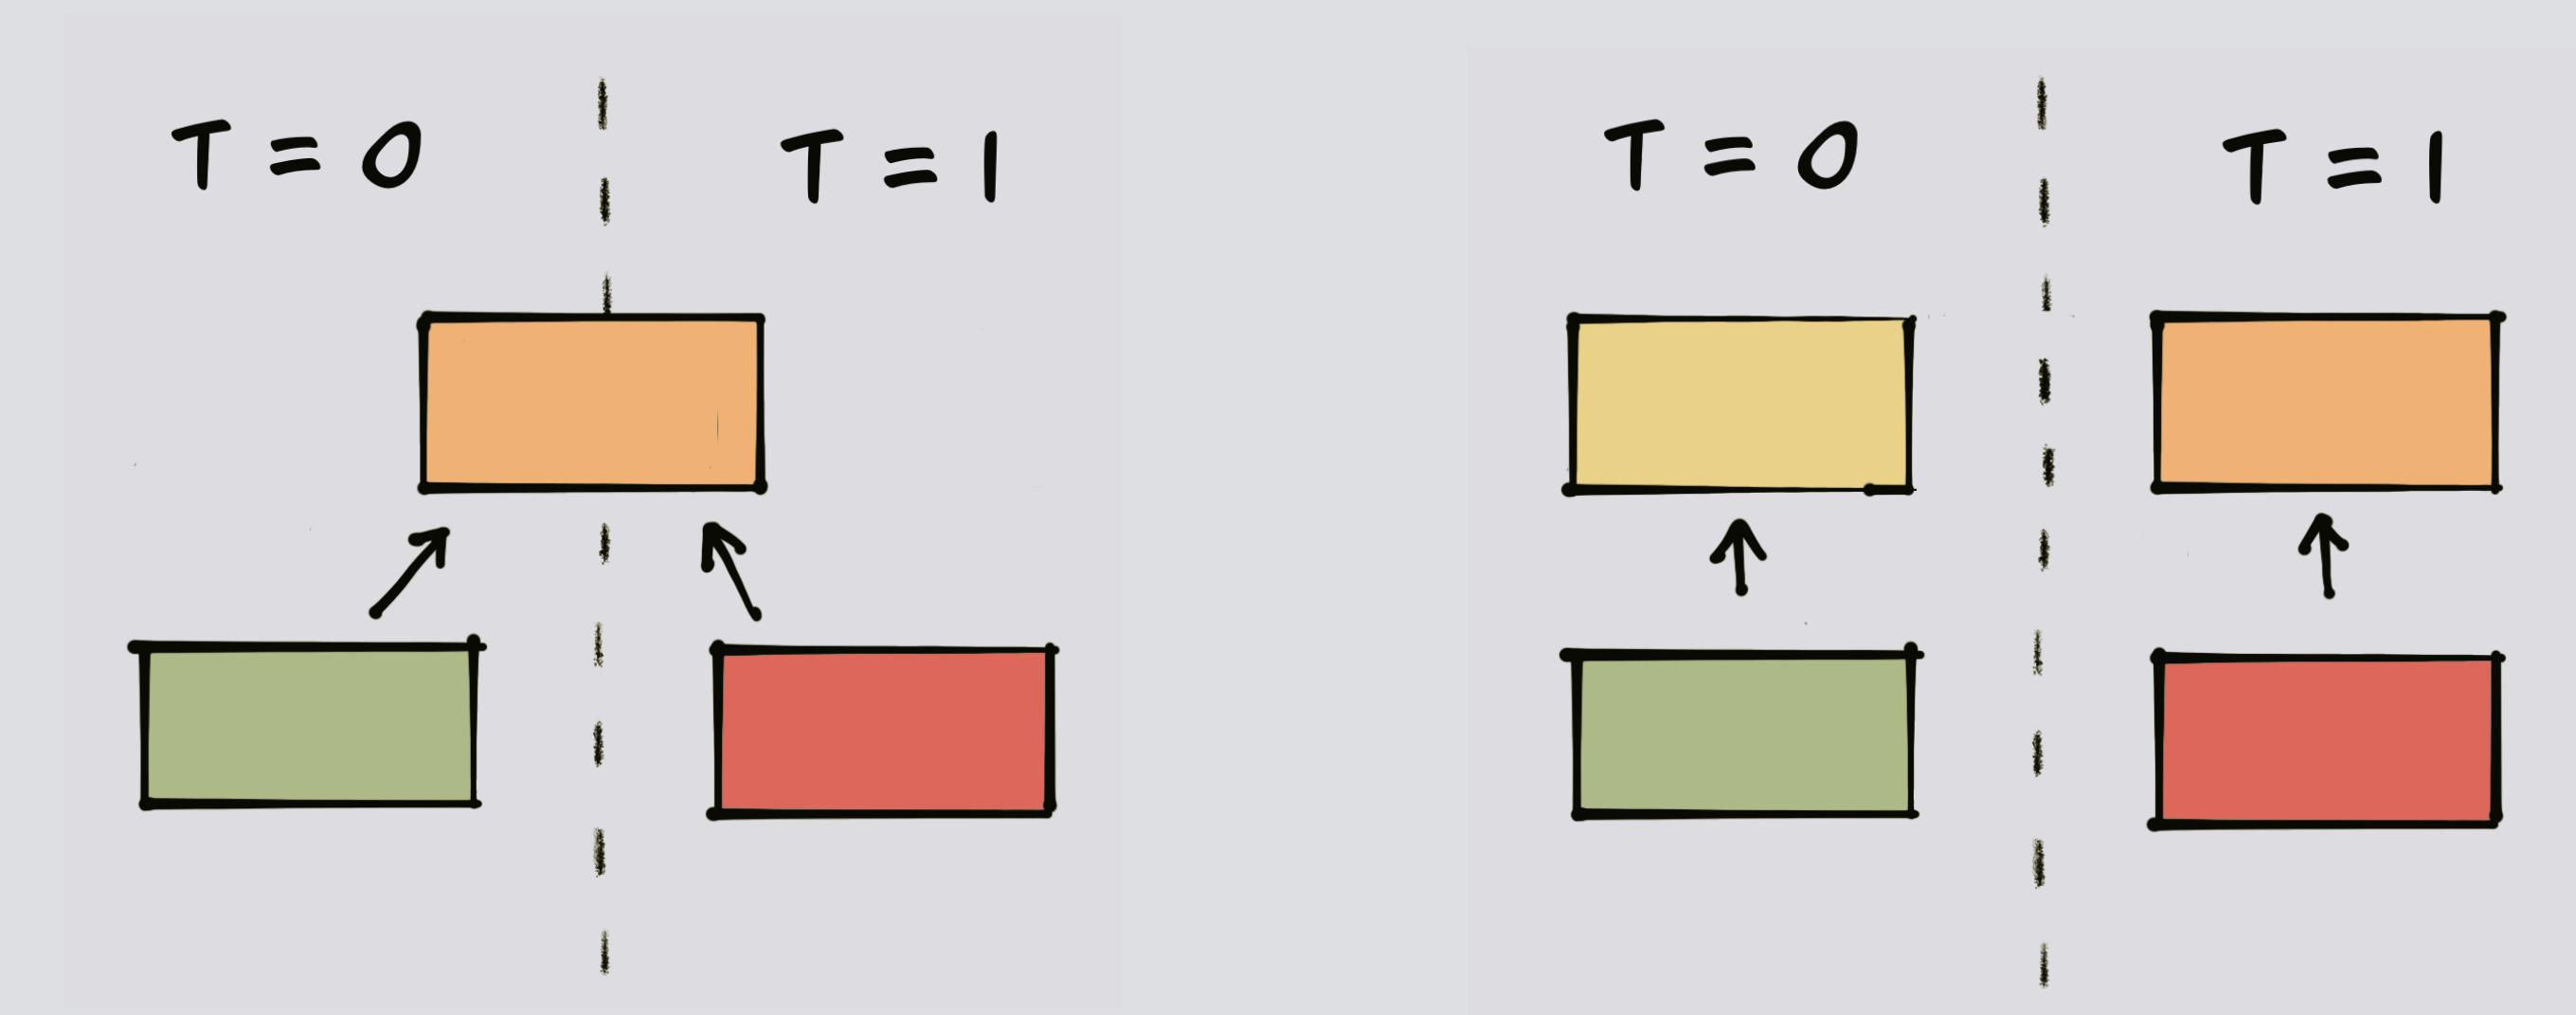
func.func @predict_dispatch_00 {
    ...
    scf.for %arg0 = %c0 to %c1024 step %c16 {
        ...
        memref.copy %19, %1 : memref<32x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>> to memref<32x16xf32, 3>
        memref.copy %20, %0 : memref<16x32xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1)>> to memref<16x32xf32, 3>
        gpu.barrier
        %21 = gpu.thread_id_x
        %22 = gpu.thread_id_y
        %23 = affine.apply affine_map<0[s0] -> (s0 * 16)>0[%22]
        %24 = affine.apply affine_map<(d0) -> ((d0 floordiv 32) * 16)>(%21)
        %25 = memref.subview %1[%23, 0] [16, 16] [1, 1] : memref<32x16xf32, 3> to
            memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 16 + s0 + d1), 3>
        %26 = memref.subview %0[0, %24] [16, 16] [1, 1] : memref<16x32xf32, 3> to
            memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1), 3>
        ...
        linalg.matmul ins(%25, %26 : memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 16 + s0 + d1), 3>,
                      memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1), 3>)
                      outs(%27 : memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>)
    }
    ...
}

```



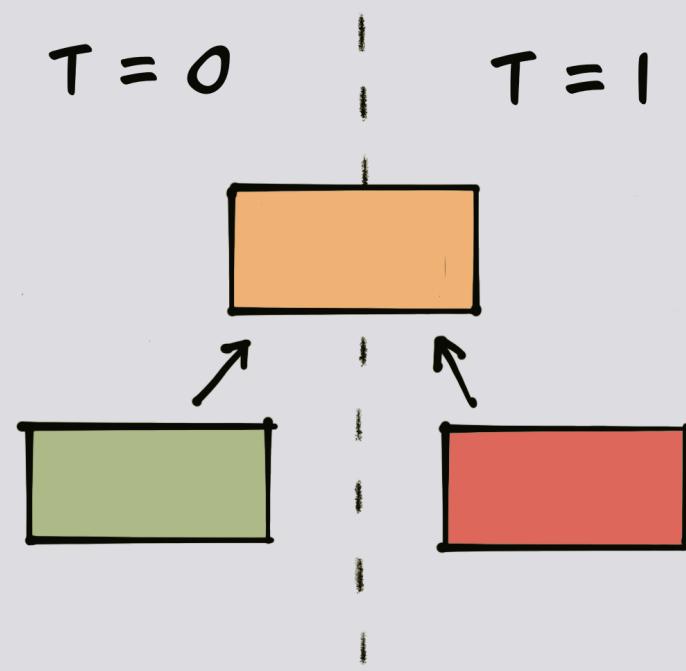
Multi-Buffering

- In order to hide latency, we can use double/multi-buffering to break dependencies between consecutive iterations of a loop using the same temporary buffer
- Required for pipelining
- Number of copies determined by desired stages of pipeline

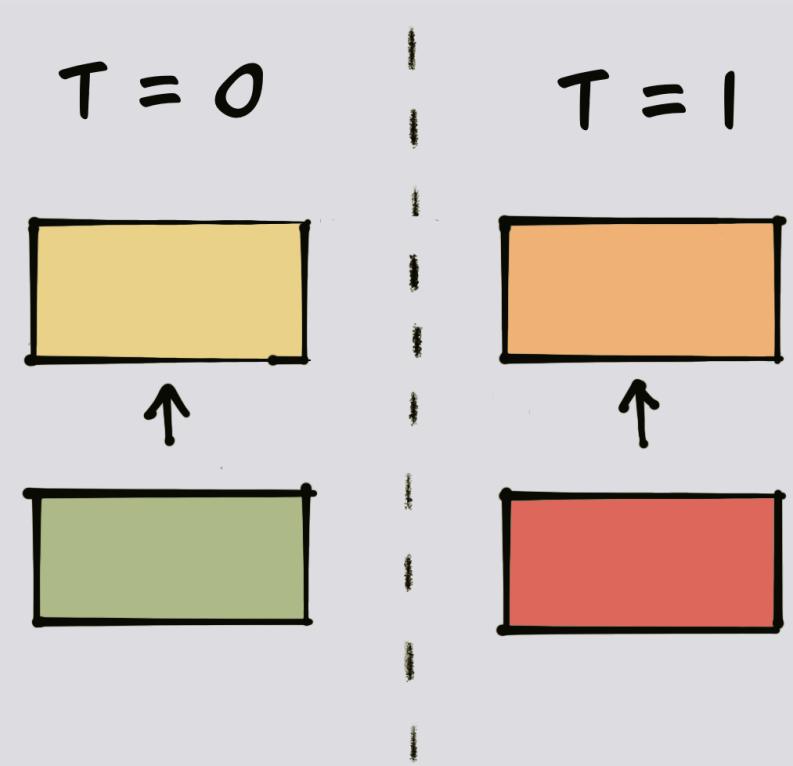


Multi-Buffering

```
func.func @predict_dispatch_00 {
    ...
    %0 = memref.alloc0 : memref<16x32xf32, 3>
    ...
    memref.copy %20, %0 : memref<16x32xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1)>> to
        memref<16x32xf32, 3>
}
```



```
func.func @predict_dispatch_00 {
    ...
    %0 = memref.alloc0 : memref<4x16x32xf32, 3>
    ...
    %22 = memref.subview %0[%21, 0, 0] [1, 16, 32] [1, 1, 1] : memref<4x16x32xf32, 3> to
        memref<16x32xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1), 3>
    ...
    memref.copy %23, %20 : memref<32x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>> to
        memref<32x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 16 + s0 + d1), 3>
}
```



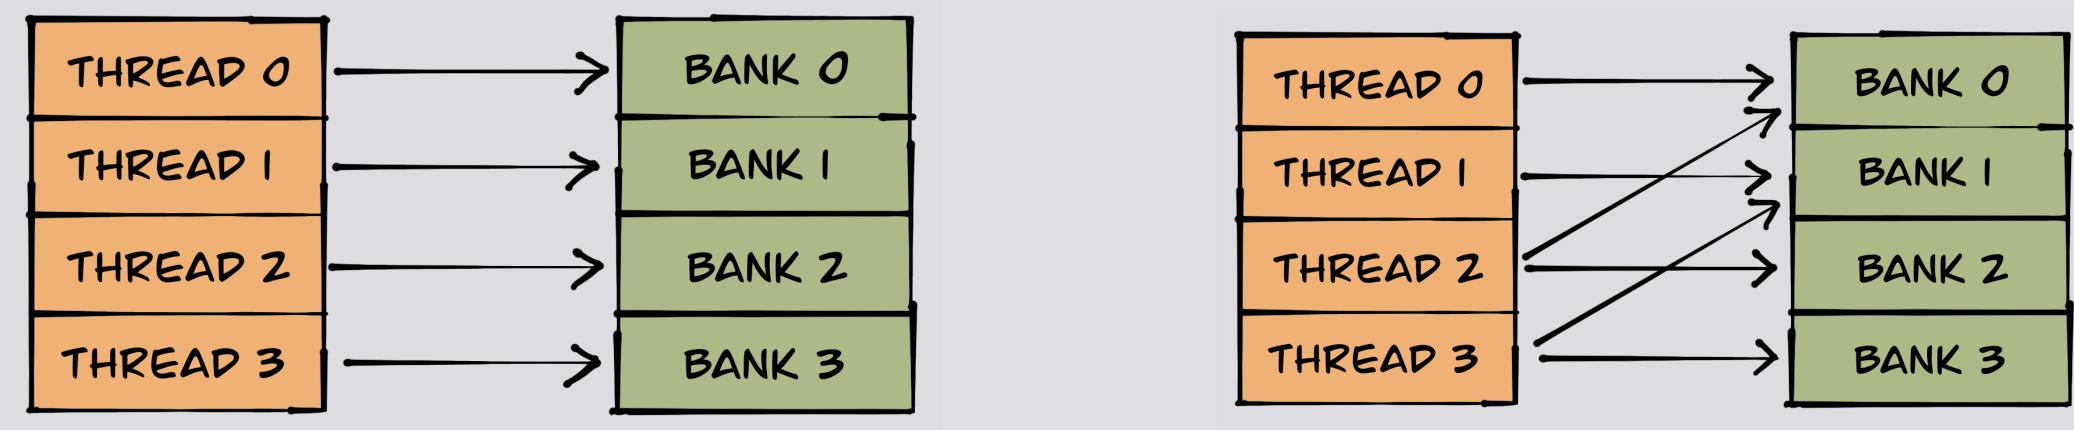
Vectorize Shared Memory Copies

- Vectorize the shared memory copy (converts to `vector.transfer_read`, `vector.transfer_write`)
- For optimal performance, always want to copy 128 bits. This can be used to determine copy tile size.
- Unroll the vector transfer read and writes

```
%3 = memref.alloc0 : memref<4x16x32xf32, 3>
%4 = memref.alloc0 : memref<4x32x16xf32, 3>
...
linalg.fill ins(%cst : f32) outs(%14 : memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>)
scf.for %arg0 = %c0 to %c1024 step %c16 {
    ...
    %22 = vector.transfer_read %18[%20, %21], %cst {in_bounds = [true, true]} :
        memref<32x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>, vector<1x4xf32>
    vector.transfer_write %22, %16[%20, %21] {in_bounds = [true, true]} : vector<1x4xf32>,
        memref<32x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 16 + s0 + d1)>, 3>
    ...
    %25 = vector.transfer_read %19[%23, %24], %cst {in_bounds = [true, true]} :
        memref<16x32xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1)>>, vector<1x4xf32>
    vector.transfer_write %25, %17[%23, %24] {in_bounds = [true, true]} : vector<1x4xf32>,
        memref<16x32xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1)>, 3>
    gpu.barrier
    ...
    linalg.matmul ins(%27, %28 : memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 16 + s0 + d1)>, 3>,
                    memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 32 + s0 + d1)>, 3>)
                    outs(%29 : memref<16x16xf32, ...>)
    }
    linalg.generic {indexing_maps = [affine_map<(d0, d1) -> (d0, d1)>],
                    iterator_types = ["parallel", "parallel"]}
                    outs(%14 : memref<16x16xf32, affine_map<(d0, d1)[s0] -> (d0 * 1024 + s0 + d1)>>) {
    ...
}
```

Reduce Bank Conflicts

- Shared memory is arranged in banks (32 banks each of width 4 bytes)
- Each thread in a warp can access shared memory in parallel
- When 2 or more threads in a warp access 4 byte words in the same bank, results in serialized accesses and hence a reduction in overall bandwidth
- Pad inner dimensions of allocOp to reduce chances of having bank conflicts (with 16 bytes)
- Plan to switch to shared memory swizzle for better efficiency



```
func.func @predict_dispatch_00 {
    ...
    %3 = memref.alloc0 : memref<4x16x32xf32, 3>
    %4 = memref.subview %3[0, 0, 0] [4, 16, 32] [1, 1, 1] : memref<4x16x36xf32, 3> to ...
    %5 = memref.alloc0 : memref<4x32x16xf32, 3>
    %6 = memref.subview %5[0, 0, 0] [4, 32, 16] [1, 1, 1] : memref<4x32x20xf32, 3> to ...
    ...
}
```

```
func.func @predict_dispatch_00 {
    ...
    %3 = memref.alloc0 : memref<4x16x36xf32, 3>
    %4 = memref.subview %3[0, 0, 0] [4, 16, 32] [1, 1, 1] : memref<4x16x36xf32, 3> to ...
    %5 = memref.alloc0 : memref<4x32x20xf32, 3>
    %6 = memref.subview %5[0, 0, 0] [4, 32, 16] [1, 1, 1] : memref<4x32x20xf32, 3> to ...
    ...
}
```

Tensor Core Vectorization

- Replace all linalg ops with vector equivalents
- Lower linalg.matmul to vector.contract with vector.transfer_read/write
- Lower linalg.generic to arith ops with vector.transfer_read/write
- Optimize vector transfers by removing redundant ops
- Vector transfer/contract size determined by tensor core supported sizes

```
%16 = scf.for %arg0 = %c0 to %c1024 step %c16 iter_args(%arg1 = %cst) -> (vector<16x16xf32>) {
    gpu.barrier
    %25 = vector.transfer_read %6[%24, %23], %cst_0 : memref<64x1024xf32>, vector<4xf32>
    vector.transfer_write %25, %4[%26, %12, %13] : vector<4xf32>, memref<4x32x20xf32, 3>
    %28 = vector.transfer_read %5[%27, %15], %cst_0 : memref<1024x32xf32>, vector<4xf32>
    vector.transfer_write %28, %3[%29, %14, %15] : vector<4xf32>, memref<4x16x36xf32, 3>
    gpu.barrier
    ...
    %32 = vector.transfer_read %4[%31, %30, %c0], %cst_0 {in_bounds = [true, true]} :
        memref<4x32x20xf32, 3>, vector<16x8xf32>
    ...
    %41 = vector.transfer_read %3[%40, %c8, %39], %cst_0 {in_bounds = [true, true]} :
        memref<4x16x36xf32, 3>, vector<8x16xf32>
    %42 = vector.contract {indexing_maps = [affine_map<(d0, d1, d2) -> (d0, d2)>,
        affine_map<(d0, d1, d2) -> (d2, d1)>,
        affine_map<(d0, d1, d2) -> (d0, d1)>],
        iterator_types = ["parallel", "parallel", "reduction"], kind = #vector.kind<add>}
        %32, %38, %arg1 : vector<16x8xf32>, vector<8x16xf32> into vector<16x16xf32>
    %43 = vector.contract {indexing_maps = [affine_map<(d0, d1, d2) -> (d0, d2)>,
        affine_map<(d0, d1, d2) -> (d2, d1)>,
        affine_map<(d0, d1, d2) -> (d0, d1)>],
        iterator_types = ["parallel", "parallel", "reduction"], kind = #vector.kind<add>}
        %35, %41, %42 : vector<16x8xf32>, vector<8x16xf32> into vector<16x16xf32>
    scf.yield %43 : vector<16x16xf32>
}
%17 = arith.maxf %16, %cst : vector<16x16xf32>
```

Convert to GPU Dialect

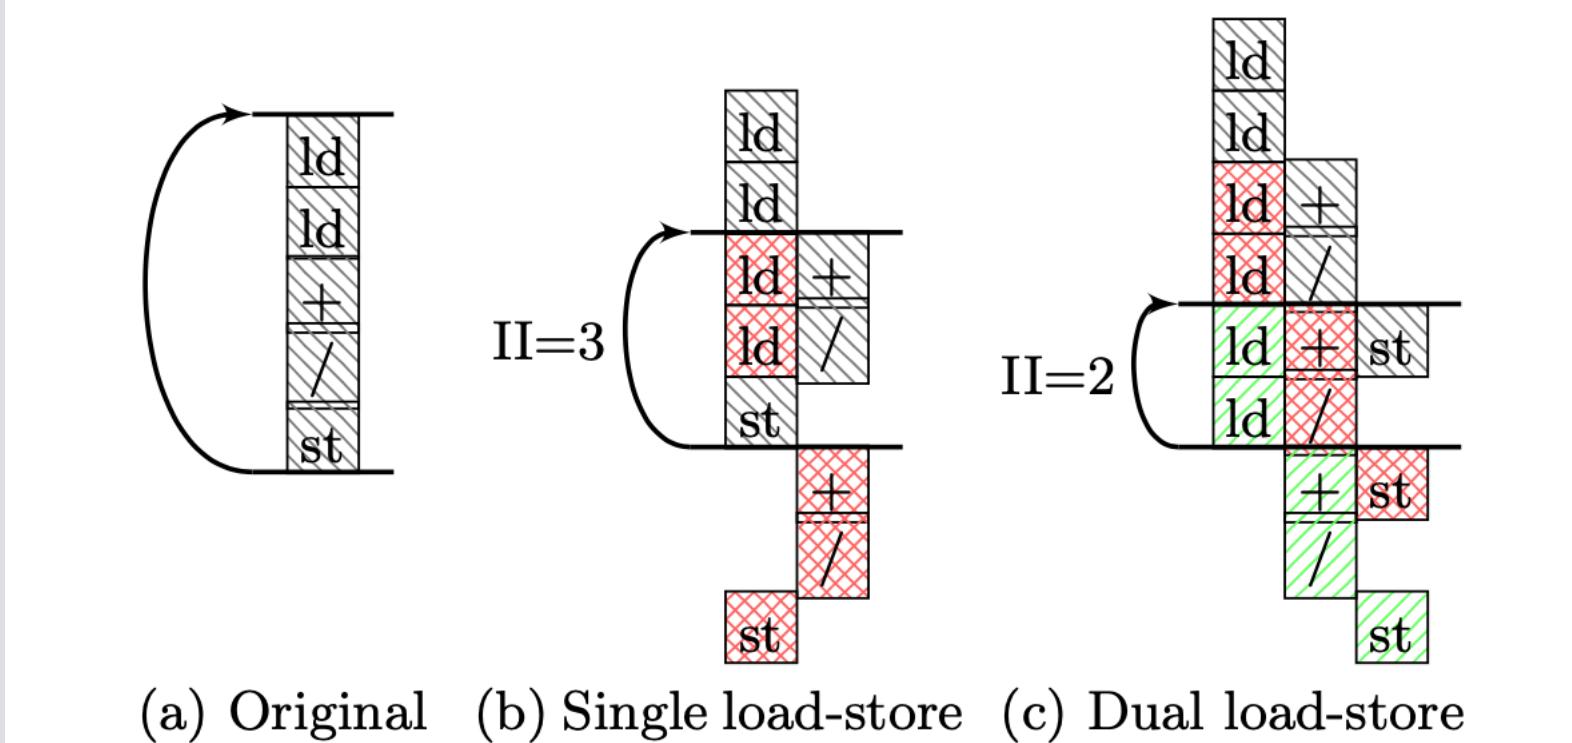
- Convert copies to shared memory to async copies
(nvgpu.device_async_copy,
nvgpu.device_async_create_group,
nvgpu.device_async_wait)
- Adds gpu.subgroup_mma_load_matrix which loads a matrix using all threads in a subgroup
- Lowers vector.contract to
gpu.subgroup_mma_compute which performs matrix-multiply accumulate using all threads in the subgroup
- Lowers arith elementwise ops to
gpu.subgroup_mma_elementwise

```
%13 = scf.for %arg0 = %c0 to %c1024 step %c16 iter_args(%arg1 = %0) ->
  (!gpu.mma_matrix<16x16xf32, "COp") {
    gpu.barrier
    %20 = nvgpu.device_async_copy %7[%18, %17], %5[%19, %9, %10], 4 : memref<64x1024xf32>
      to memref<4x32x20xf32, 3>
    %22 = nvgpu.device_async_copy %6[%21, %12], %4[%19, %11, %12], 4 : memref<1024x32xf32>
      to memref<4x16x36xf32, 3>
    %23 = nvgpu.device_async_create_group %20, %22
    nvgpu.device_async_wait %23
    gpu.barrier
    ...
    %25 = gpu.subgroup_mma_load_matrix %5[%19, %24, %c0] {leadDimension = 20 : index} :
      memref<4x32x20xf32, 3> -> !gpu.mma_matrix<16x8xf32, "AOp">
    %28 = gpu.subgroup_mma_load_matrix %4[%19, %c0, %27] {leadDimension = 36 : index} :
      memref<4x16x36xf32, 3> -> !gpu.mma_matrix<8x16xf32, "BOp">
    %31 = gpu.subgroup_mma_compute %26, %29, %30 :
      !gpu.mma_matrix<16x8xf32, "AOp">, !gpu.mma_matrix<8x16xf32, "BOp">
      -> !gpu.mma_matrix<16x16xf32, "COp">
      scf.yield %31 : !gpu.mma_matrix<16x16xf32, "COp">
}
%14 = gpu.subgroup_mma_elementwise maxf %13, %0 :
  (!gpu.mma_matrix<16x16xf32, "COp">, !gpu.mma_matrix<16x16xf32, "COp">) ->
  !gpu.mma_matrix<16x16xf32, "COp">
  gpu.subgroup_mma_store_matrix %14, %8[%15, %16] {leadDimension = 1024 : index} :
  !gpu.mma_matrix<16x16xf32, "COp">, memref<64x1024xf32>
```

GPU Pipelining

- GPU Pipelining
 - Implement software pipelining using modulo scheduling
 - Operations of original loop body are overlapped so that there is a fixed initiation interval (II) between the start of consecutive loop iterations
 - Scheduling can be constrained by available hardware resource or loop carried dependencies
 - Emit a prologue, kernel and epilogue

```
for (int i = 0; i < N; i++) {
    B[i] = (A[i] + A[i + 1]) / 4;
}
```



GPU Pipelining

- Prologue contains async copies to shared memory
- Kernel contains mma load matrix and compute and ends with async copies
- Epilogue contains remaining computation

↑ Prologue → ↓ Kernel ← Epilogue ↑

```

gpu.barrier __pipeline_global_load__
%14 = nvgpu.device_async_copy %7[%13, %10], %5[%c0, %9, %10], 4 __pipeline_global_load__ : memref<64x1024xf32> to memref<4x32x20xf32, 3>
%15 = nvgpu.device_async_copy %6[%11, %12], %4[%c0, %11, %12], 4 __pipeline_global_load__ : memref<1024x32xf32> to memref<4x16x36xf32, 3>
%16 = nvgpu.device_async_create_group %14, %15 __pipeline_global_load__

...
gpu.barrier __pipeline_global_load__
%28 = nvgpu.device_async_copy %7[%13, %27], %5[%c3, %9, %10], 4 __pipeline_global_load__ : memref<64x1024xf32> to memref<4x32x20xf32, 3>
%30 = nvgpu.device_async_copy %6[%29, %12], %4[%c3, %11, %12], 4 __pipeline_global_load__ : memref<1024x32xf32> to memref<4x16x36xf32, 3>
%31 = nvgpu.device_async_create_group %28, %30 __pipeline_global_load__
%32:9 = scf.for %arg0 = %c0 to %c960 step %c16 iter_args%arg1 = %0, %arg2 = %16, %arg3 = %21, %arg4 = %26, %arg5 = %31, %arg6 = %c0, %arg7 = %c1, %arg8 = %c2, %arg9 = %c3
    nvgpu.device_async_wait %arg2 {numGroups = 3 : i32}
    gpu.barrier
    %63 = gpu.subgroup_mma_load_matrix %5[%arg6, %62, %c0] {leadDimension = 20 : index} : memref<4x32x20xf32, 3> -> lgpu.mma_matrix<16x8xf32, "AOp">
    %66 = gpu.subgroup_mma_load_matrix %4[%arg6, %c0, %65] {leadDimension = 36 : index} : memref<4x16x36xf32, 3> -> lgpu.mma_matrix<8x16xf32, "BOp">
    %68 = gpu.subgroup_mma_compute %63, %66, %arg1 : lgpu.mma_matrix<16x8xf32, "AOp">, lgpu.mma_matrix<8x16xf32, "BOp"> -> lgpu.mma_matrix<16x16xf32, "COp">
    gpu.barrier __pipeline_global_load__
    %70 = arith.addi %arg0, %c64 : index
    %73 = nvgpu.device_async_copy %7[%13, %71], %5[%72, %9, %10], 4 __pipeline_global_load__ : memref<64x1024xf32> to memref<4x32x20xf32, 3>
    %75 = nvgpu.device_async_copy %6[%74, %12], %4[%72, %11, %12], 4 __pipeline_global_load__ : memref<1024x32xf32> to memref<4x16x36xf32, 3>
    %76 = nvgpu.device_async_create_group %73, %75 __pipeline_global_load__
    scf.yield %69, %arg3, %arg4, %arg5, %76, %arg7, %arg8, %arg9, %72 : ...
}

nvgpu.device_async_wait %32#1 {numGroups = 3 : i32}
gpu.barrier
%34 = gpu.subgroup_mma_load_matrix %5[%32#5, %33, %c0] {leadDimension = 20 : index} : memref<4x32x20xf32, 3> -> lgpu.mma_matrix<16x8xf32, "AOp">
%37 = gpu.subgroup_mma_load_matrix %4[%32#5, %c0, %36] {leadDimension = 36 : index} : memref<4x16x36xf32, 3> -> lgpu.mma_matrix<8x16xf32, "BOp">
%39 = gpu.subgroup_mma_compute %34, %37, %32#0 : lgpu.mma_matrix<16x8xf32, "AOp">, lgpu.mma_matrix<8x16xf32, "BOp"> -> lgpu.mma_matrix<16x16xf32, "COp">
nvgpu.device_async_wait %32#2 {numGroups = 2 : i32}
gpu.barrier

...
%53 = gpu.subgroup_mma_load_matrix %5[%32#8, %33, %c0] {leadDimension = 20 : index} : memref<4x32x20xf32, 3> -> lgpu.mma_matrix<16x8xf32, "AOp">
%55 = gpu.subgroup_mma_load_matrix %4[%32#8, %c0, %36] {leadDimension = 36 : index} : memref<4x16x36xf32, 3> -> lgpu.mma_matrix<8x16xf32, "BOp">
%57 = gpu.subgroup_mma_compute %53, %55, %52 : lgpu.mma_matrix<16x8xf32, "AOp">, lgpu.mma_matrix<8x16xf32, "BOp"> -> lgpu.mma_matrix<16x16xf32, "COp">
%59 = gpu.subgroup_mma_elementwise maxf %58, %0 : (lgpu.mma_matrix<16x16xf32, "COp">, lgpu.mma_matrix<16x16xf32, "COp">) -> lgpu.mma_matrix<16x16xf32, "COp">
gpu.subgroup_mma_store_matrix %59, %8[%60, %61] {leadDimension = 1024 : index} : lgpu.mma_matrix<16x16xf32, "COp">, memref<64x1024xf32>

```

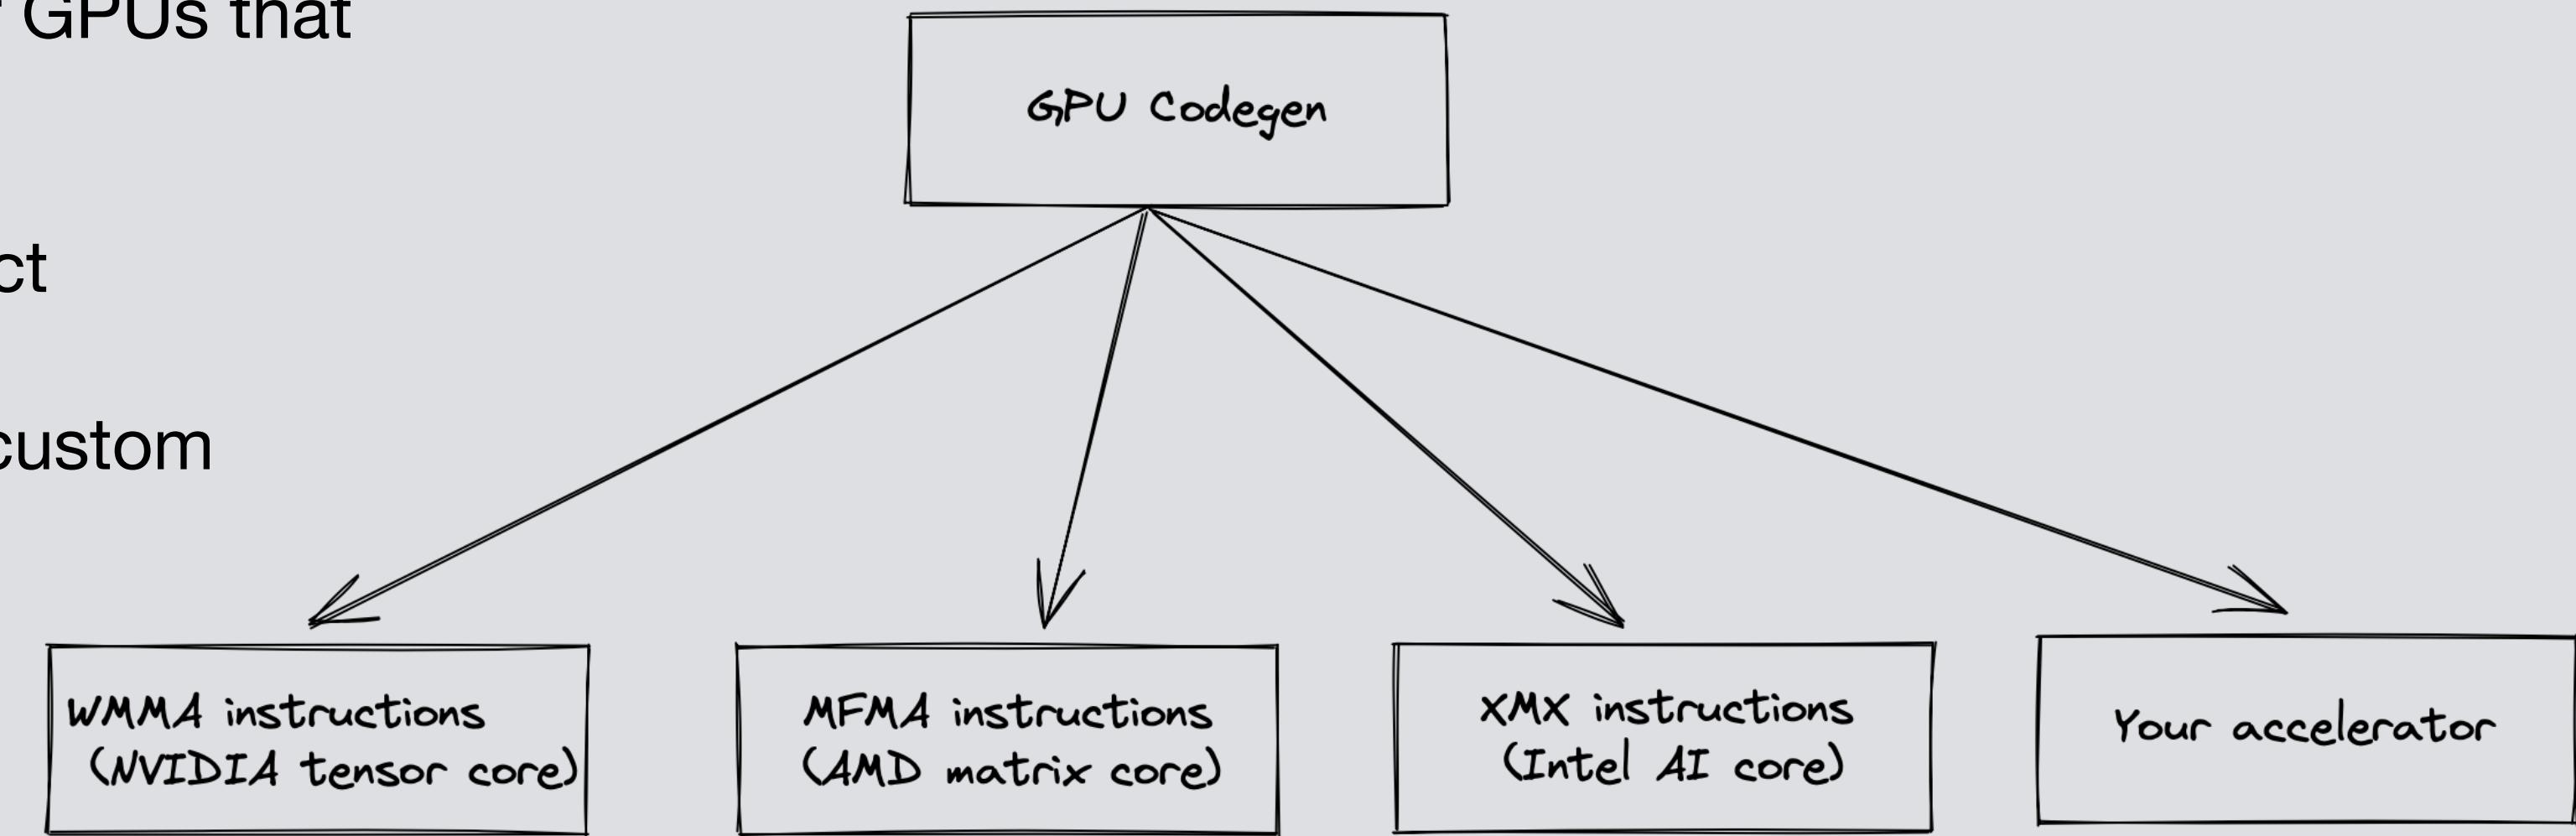
GPU Lowering

- Next, we use LLVM to lower to PTX
- Use JIT compilation to convert PTX code to native GPU machine code
- Alternatively, could use PTX assembler and convert to SASS (and then CUBIN)

```
wmma.load.a.sync.aligned.row.m16n16k8.shared.tf32 [%r43, %r44, %r45, %r46], [%rd107+11776], %r6;
wmma.load.a.sync.aligned.row.m16n16k8.shared.tf32 [%r47, %r48, %r49, %r50], [%rd107+11808], %r6;
wmma.load.b.sync.aligned.row.m16n16k8.shared.tf32 [%r51, %r52, %r53, %r54], [%rd109+2304], %r15;
wmma.load.b.sync.aligned.row.m16n16k8.shared.tf32 [%r55, %r56, %r57, %r58], [%rd109+3456], %r15;
wmma.mma.sync.aligned.row.row.m16n16k8.f32+tf32.f32
    [%f49, %f50, %f51, %f52, %f53, %f54, %f55, %f56],
    [%r43, %r44, %r45, %r46],
    [%r51, %r52, %r53, %r54],
    [%f41, %f42, %f43, %f44, %f45, %f46, %f47, %f48];
wmma.mma.sync.aligned.row.row.m16n16k8.f32+tf32.f32
    [%f57, %f58, %f59, %f60, %f61, %f62, %f63, %f64],
    [%r47, %r48, %r49, %r50],
    [%r55, %r56, %r57, %r58],
    [%f49, %f50, %f51, %f52, %f53, %f54, %f55, %f56];
cp.async.wait_group 1;
bar.sync 0;
wmma.load.a.sync.aligned.row.m16n16k8.shared.tf32 [%r59, %r60, %r61, %r62], [%rd107+14336], %r6;
wmma.load.a.sync.aligned.row.m16n16k8.shared.tf32 [%r63, %r64, %r65, %r66], [%rd107+14368], %r6;
wmma.load.b.sync.aligned.row.m16n16k8.shared.tf32 [%r67, %r68, %r69, %r70], [%rd109+4608], %r15;
wmma.load.b.sync.aligned.row.m16n16k8.shared.tf32 [%r71, %r72, %r73, %r74], [%rd109+5760], %r15;
wmma.mma.sync.aligned.row.row.m16n16k8.f32+tf32.f32
    [%f65, %f66, %f67, %f68, %f69, %f70, %f71, %f72],
    [%r59, %r60, %r61, %r62],
    [%r67, %r68, %r69, %r70],
    [%f57, %f58, %f59, %f60, %f61, %f62, %f63, %f64];
wmma.mma.sync.aligned.row.row.m16n16k8.f32+tf32.f32
    [%f73, %f74, %f75, %f76, %f77, %f78, %f79, %f80],
    [%r63, %r64, %r65, %r66],
    [%r71, %r72, %r73, %r74],
    [%f65, %f66, %f67, %f68, %f69, %f70, %f71, %f72];
cp.async.wait_group 0;
bar.sync 0;
```

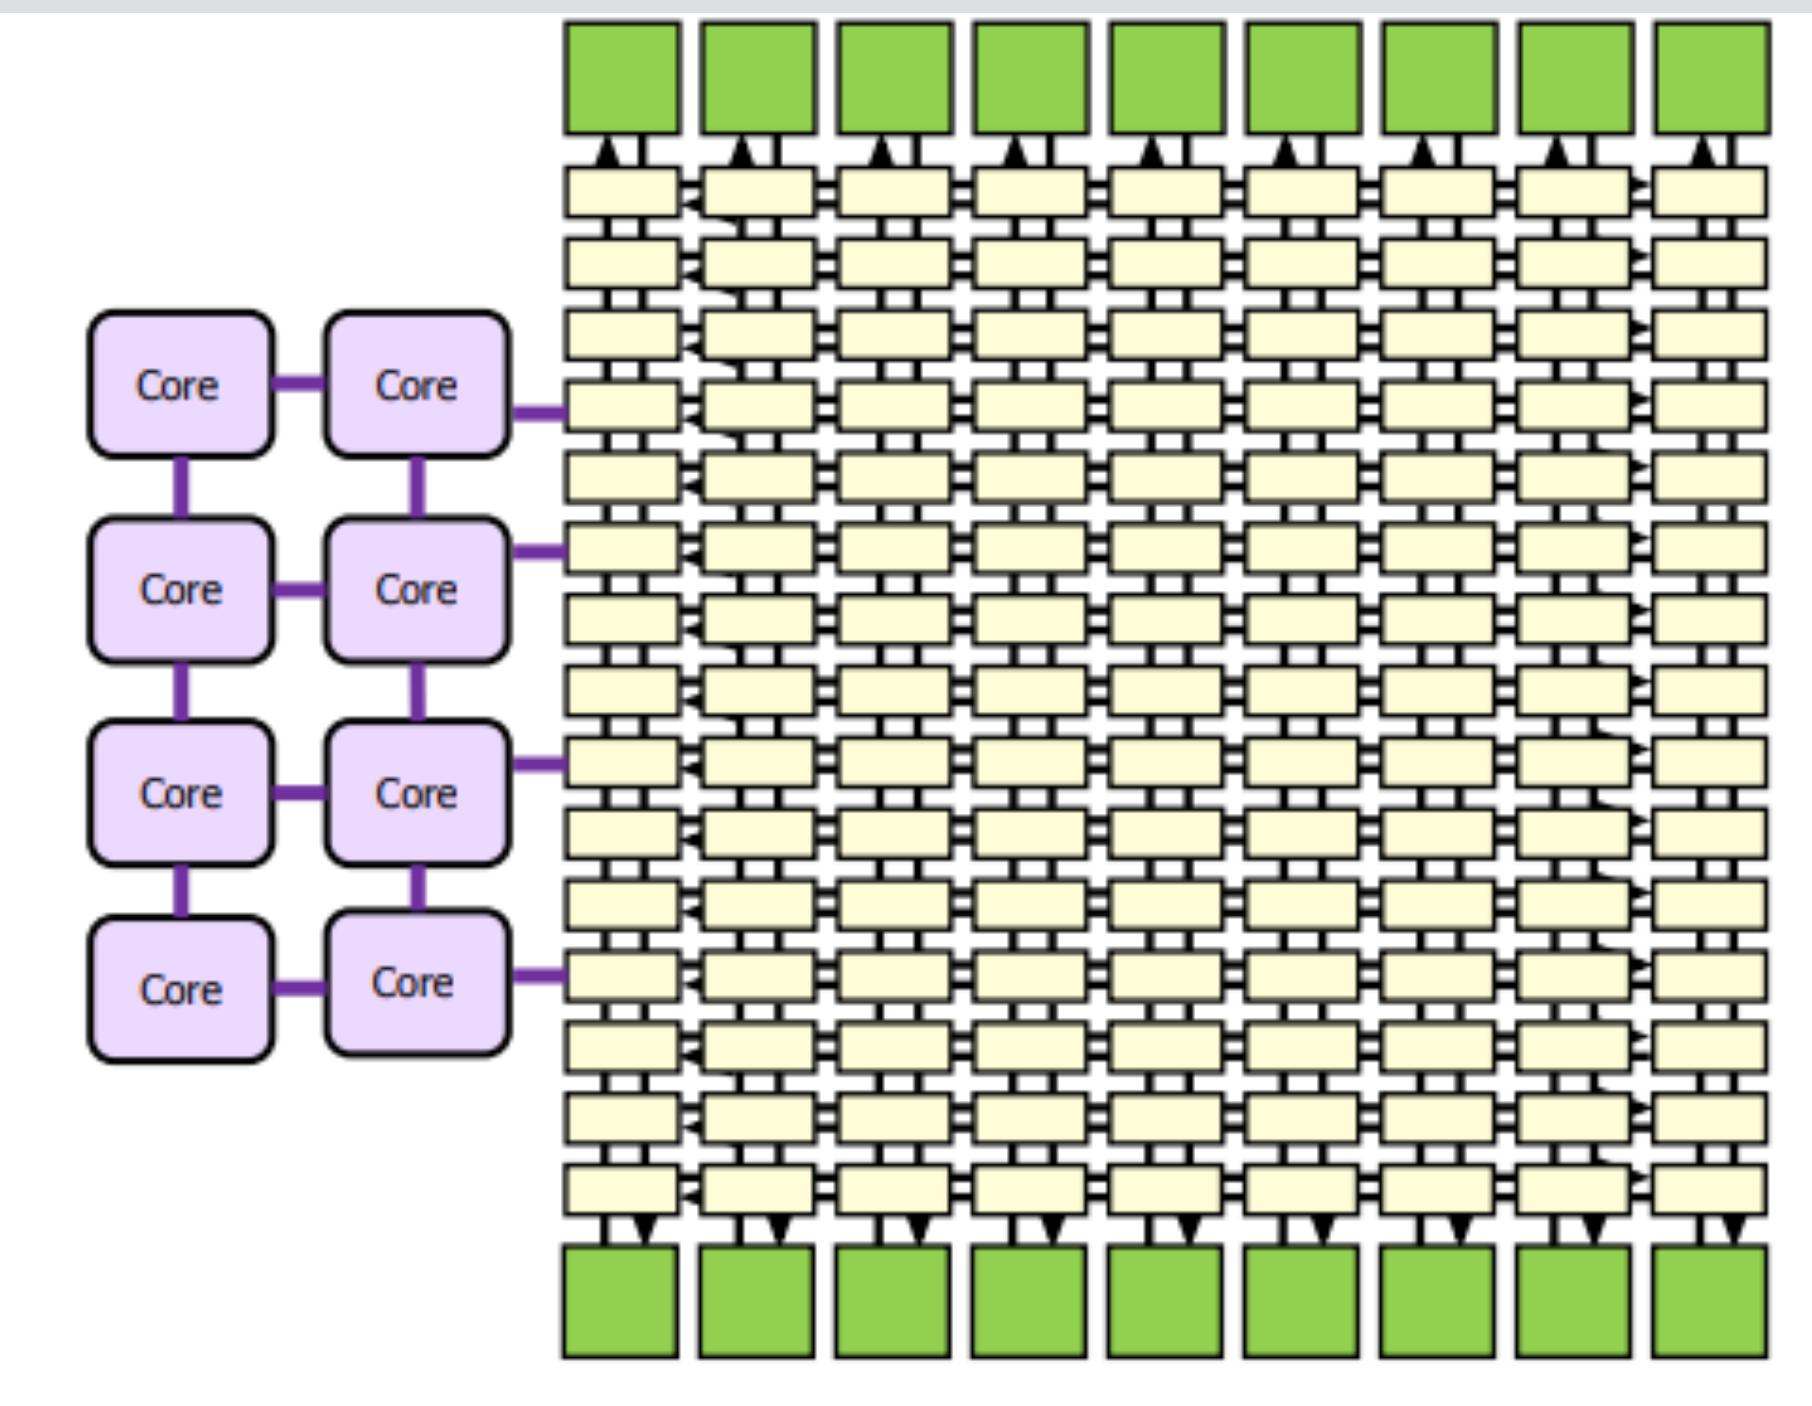
Additional Targets

- Can reuse existing pipeline to target other GPUs that have tensor core equivalent functionality
- Adjust lowering from linalg to vector dialect
- Add appropriate dialect and lowering for custom accelerator



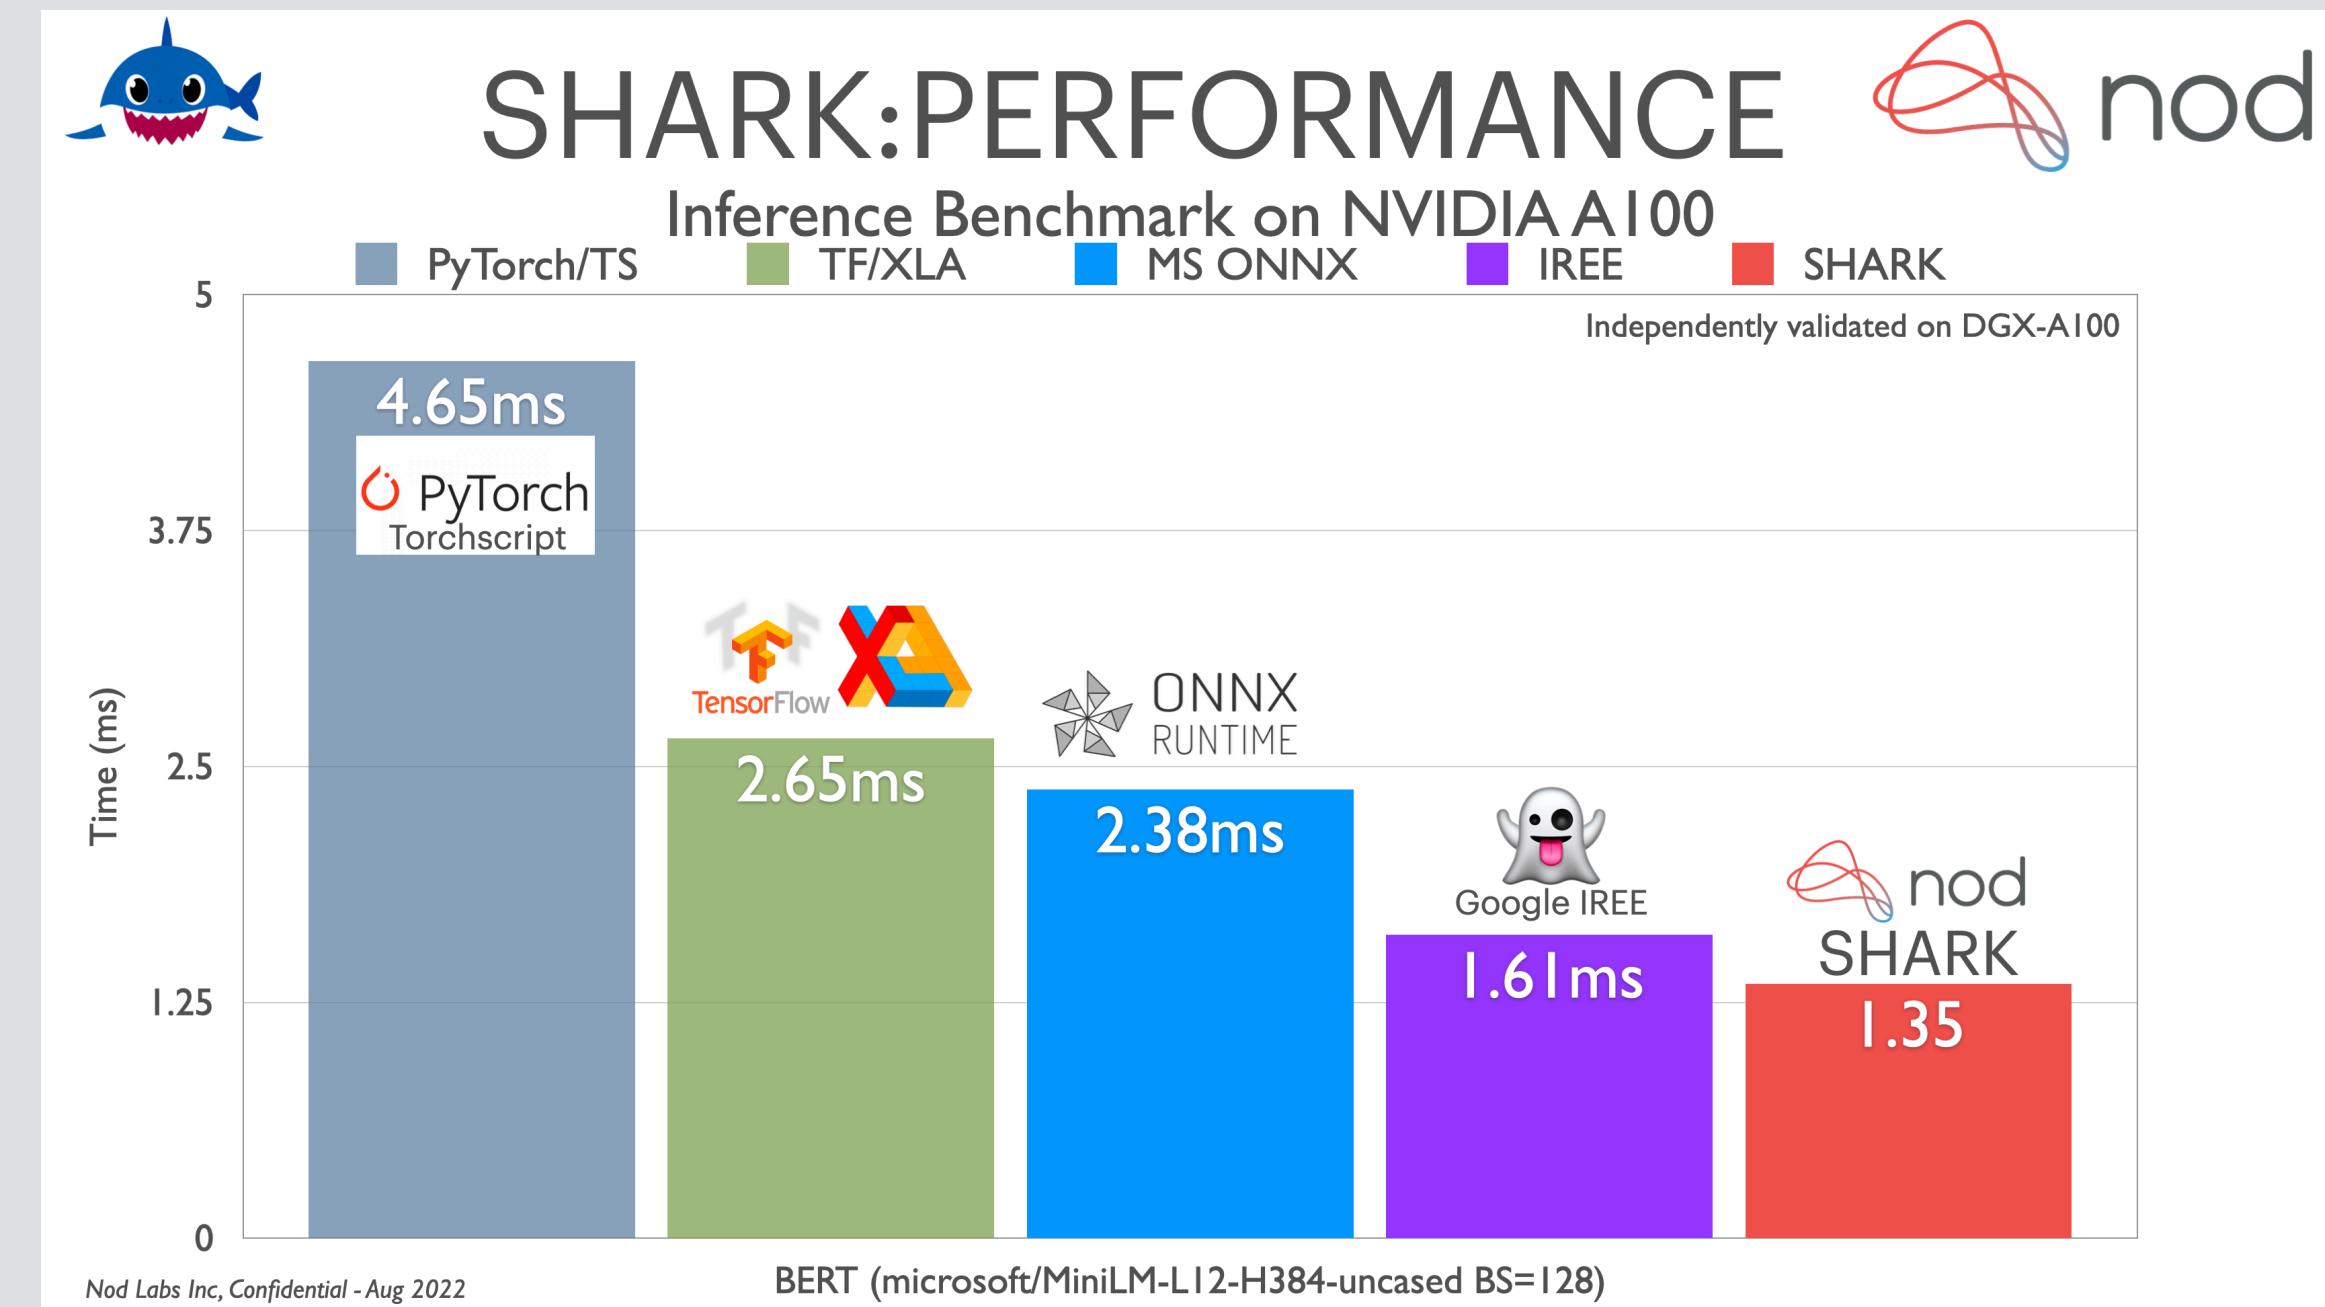
Targeting custom accelerators

- New RISC-V based many-core architectures (such as Hammerblade)
- Create a new dialect to model multiple processing elements (PE) and memory hierarchy
- Can leverage vector dialect or experimental RVV dialect for vector code generation
- Need to develop cost model to determine how to place kernels on tiles of PEs
- Could have special function units (tensor core equivalents) that can be leveraged during code generation for better performance



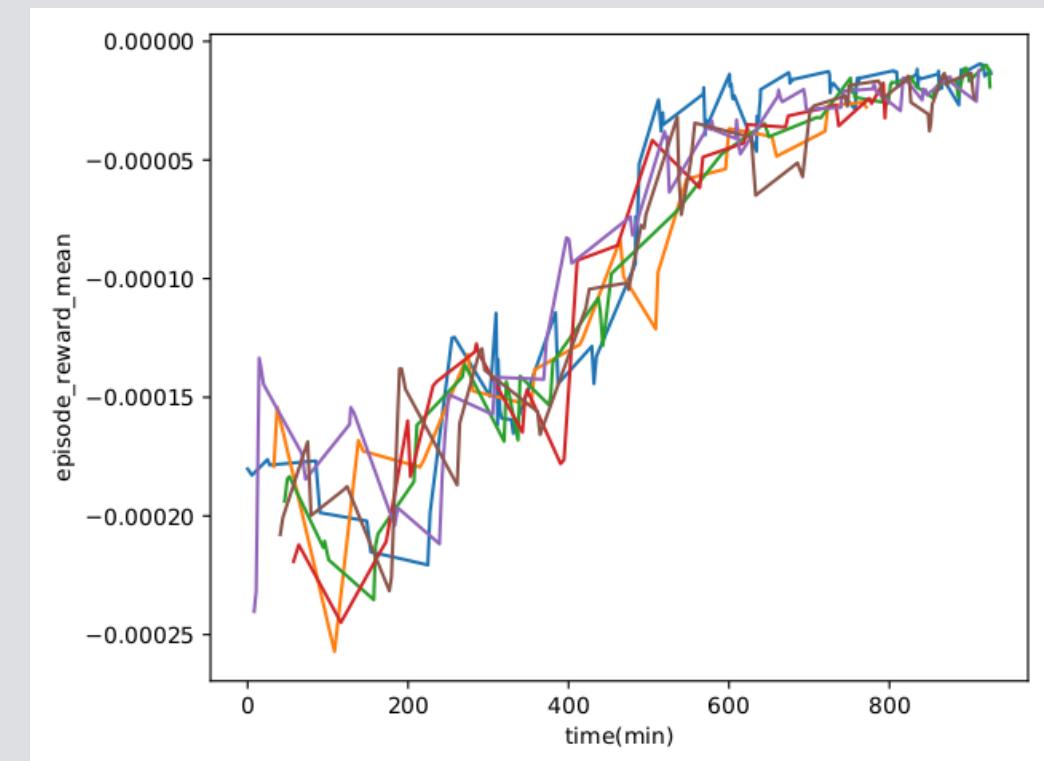
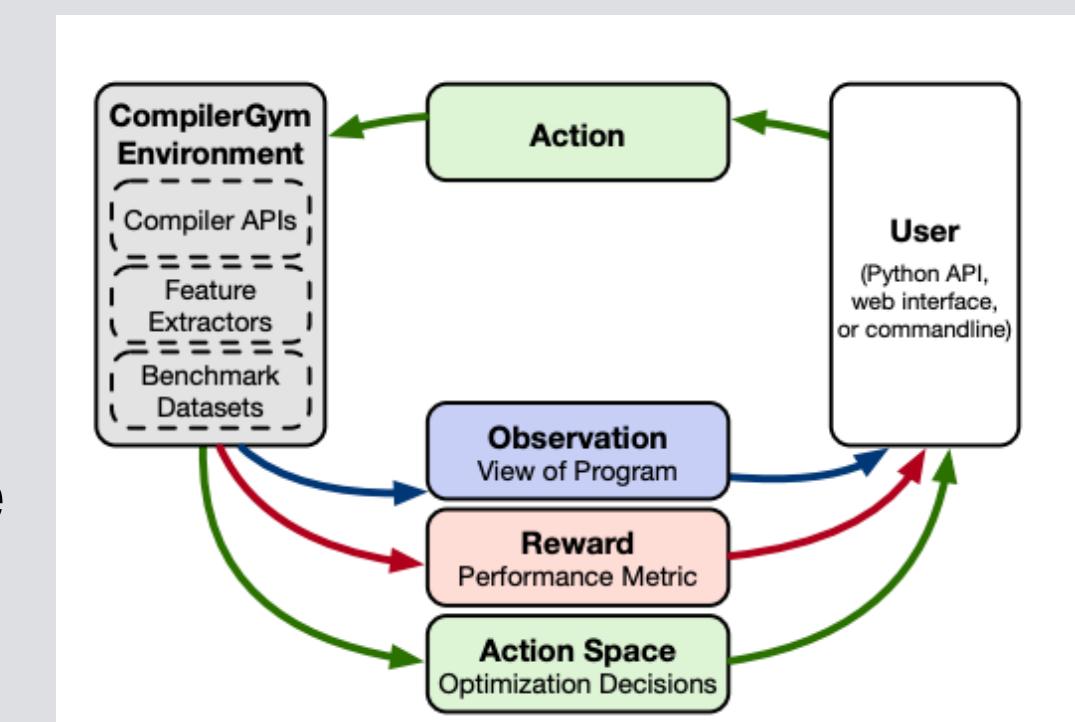
Auto-tuning

- How to determine optimal tile sizes?
- What about loop interchange? How much to unroll the loop?
- How does this extend from a single operator to entire neural network? On distributed heterogeneous resources?
- Many hyperparameters and a large search space make it difficult to generate good code
- Unclear how much performance is left on table for a given set of hyperparameters
- Formulate as search problem imposing constraints to reduce search space
- Can choose from a variety of search algorithms ranging from reinforcement learning (RL) to genetic algorithms



Auto-tuning

- RL Framework (Compiler Gym) that can be used for compiler optimization tasks
- Gradient-free methods (Nevergrad) have also been used to tune performance
- Current list of knobs shown in table below
- Beyond operator level, we can also search for how to best partition tensors for distributed computations (for inference and training)
- Approaches such as ALPA provide framework for how to partition tensors across clusters of heterogeneous resources
- Additional variables such as checkpointing for training



Transformation	Options
TILE	sizes – array of tile sizes interchange – order of loops after tiling pad – whether to pad partial tiles pack_paddings – non-removable padding for arrays hoist_paddings – number of loops from which to hoist padding for arrays peel – loops from which to peel off partial tiles scalarize_dyn_dims – whether to emit scalar code for non-vectorizable (dynamic) dimensions
VECTORIZE	vectorize_padding – whether to vectorize pad operations
PipelineOneParentLoop	parent_loop_num – which parent loop to pipeline II – Iteration Interval read_latency – Latency of a read operation
UnrollOneParentLoop	parent_loop_num – which parent loop to unroll unroll_amount – by how many iterations to unroll the loop
UnrollOneVectorOp	source_shape – source shape of the vector op to unroll source_shape – target shape to unroll the vector op to
Bufferize	none
Sparsify	none
LowerVectors	contraction_lowering – how to lower vector contractions (outer/inner product, LLVM matrix intrinsics) multi_reduction_lowering – how to lower multidimensional reductions (inner or outer parallel) transpose_lowering – how to lower transpositions (elementwise, flat, vector shuffle, target-specific)
LowerToLLVM	none

Conclusion

- MLIR Code generation focuses on taking high-level tensor computation primitives and lowering them to LLVM IR with appropriate intrinsics
- Attempts to take guesswork out of the backend, reducing dependence on black-box optimizers such as the LLVM auto-vectorizer
- Leverages LLVM for traditional code generation
- Many abstractions shared between CPU and GPU compilation pipeline (and potentially other new accelerators)
- Tiling and vectorization key components of both pipelines
- Additional work required to manage shared memory and target tensor cores on GPU
- Auto-tuning essential to obtaining good performance from code generated kernels
- Can be extended to handle sparse tensors, non-structured ops (`linalg.ext`)

Acknowledgements

- Nod.ai Team
 - Discord: <https://discord.gg/RUqY2h2s9u>
- Google IREE Team
 - Discord: <https://discord.gg/26P4xW4>

References

- Vasilache, N. et al. (2022) Composable and Modular Code Generation in MLIR. <https://arxiv.org/pdf/2202.03293.pdf>
- Hsin-I, C. L. et al. (2022) TinyIREE: An ML Execution Environment for Embedded Systems from Compilation to Deployment. <https://arxiv.org/pdf/2205.14479.pdf>
- Bradbury, A. LLVM backend development by example (RISC-V). <https://youtu.be/AFalP-dF-RA>
- Codegen Dialect Overview. <https://discourse.llvm.org/t/codegen-dialect-overview/2723>
- Dawkins, Q. Updated MLIR Dialect Overview Diagram. <https://discourse.llvm.org/t/rfc-updated-mlir-dialect-overview-diagram/64266>
- Linalg Dialect Rationale: The Case for Compiler Friendly Custom Operations. <https://mlir.llvm.org/docs/Rationale/RationaleLinalgDialect/>
- Anatomy of Linalg.generic. <https://youtu.be/A805W2KSCxQ>
- Loop double-buffering/multi-buffering. <https://discourse.llvm.org/t/loop-double-buffering-multi-buffering/59979>
- CUTLASS: Cuda Template Library for dense linear algebra at all levels and scales. <https://on-demand.gputechconf.com/gtc/2018/presentation/s8854-cutlass-software-primitives-for-dense-linear-algebra-at-all-levels-and-scales-within-cuda.pdf>
- Bohm, et al. A Novel Hilbert Curve for Cache-locality Preserving Loops. <https://eprints.cs.univie.ac.at/5726/1/loops.pdf>
- Jordans, R. et al. (2015) High-level software-pipelining in LLVM.
- Hammerblade Manycore Technical Reference Manual. https://docs.google.com/document/d/1b2g2nnMYidMkcn6iHJ9NGjpQYfZeWEmMdLeO_3nLtgo/edit
- Add RISC-V Vector Extension (RVV) Dialect. <https://discourse.llvm.org/t/rfc-add-risc-v-vector-extension-rvv-dialect/4146>
- Cummins, C. et al. (2021) CompilerGym: Robust, Performant Compiler Optimization Environments for AI Research. <https://arxiv.org/pdf/2109.08267.pdf>
- Zheng, L. et al (2022). Alpa: Automating Inter- and Intra-Operator Parallelism for Distributed Deep Learning. <https://arxiv.org/pdf/2201.12023.pdf>
- Images on Motivation slide taken from OpenAI, NVIDIA, AMD, Sambanova and Cerebras websites.